

```
In [1]: # import libraries
import pandas as pd
import numpy as np
import seaborn as sb
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Amazon Employee Access Challenge

Overview

When an employee at any company starts work, they first need to obtain the computer access necessary to fulfill their role. This access may allow an employee to read/manipulate resources through various applications or web portals. It is assumed that employees fulfilling the functions of a given role will access the same or similar resources. It is often the case that employees figure out the access they need as they encounter roadblocks during their daily work (e.g. not able to log into a reporting portal). A knowledgeable supervisor then takes time to manually grant the needed access in order to overcome access obstacles. As employees move throughout a company, this access discovery/recovery cycle wastes a nontrivial amount of time and money.

There is a considerable amount of data regarding an employee's role within an organization and the resources to which they have access. Given the data related to current employees and their provisioned access, models can be built that automatically determine access privileges as employees enter and leave roles within a company. These auto-access models seek to minimize the human involvement required to grant or revoke employee access.

Objective

The objective of this competition is to build a model, learned using historical data, that will determine an employee's access needs, such that manual access transactions (grants and revokes) are minimized as the employee's attributes change over time. The model will take an employee's role information and a resource code and will return whether or not access should be granted.

ML Problem

So our aim is to develop a Machine Learning model that takes an employee's access request as input which contains details about the employee's attributes like role, department etc.. and the model has to decide whether to provide access or not. Here the dataset provided by Amazon contains real historic data collected from 2010 and 2011. The Performance metric used in this case study is AUC score.

Data Information

<https://www.kaggle.com/c/amazon-employee-access-challenge/data>

Data Description

The data consists of real historical data collected from 2010 & 2011. Employees are manually allowed or denied access to resources over time. You must create an algorithm capable of learning from this historical data to predict approval/denial for an unseen set of employees.

File Descriptions

train.csv - The training set. Each row has the ACTION (ground truth), RESOURCE, and information about the employee's role at the time of approval

test.csv - The test set for which predictions should be made. Each row asks whether an employee having the listed characteristics should have access to the listed resource.

Column Descriptions

Column Name	Description
ACTION	ACTION is 1 if the resource was approved, 0 if the resource was not
RESOURCE	An ID for each resource
MGR_ID	The EMPLOYEE ID of the manager of the current EMPLOYEE ID record; an employee may have only one manager at a time
ROLE_ROLLUP_1	Company role grouping category id 1 (e.g. US Engineering)
ROLE_ROLLUP_2	Company role grouping category id 2 (e.g. US Retail)
ROLE_DEPTNAME	Company role department description (e.g. Retail)
ROLE_TITLE	Company role business title description (e.g. Senior Engineering Retail Manager)
ROLE_FAMILY_DESC	Company role family extended description (e.g. Retail Manager, Software Engineering)
ROLE_FAMILY	Company role family description (e.g. Retail Manager)
ROLE_CODE	Company role code; this code is unique to each role (e.g. Manager)

Data Analysis

```
In [2]: train = pd.read_csv('data/train.csv')
test = pd.read_csv('data/test.csv')
```

```
In [3]: train.shape
```

```
Out[3]: (32769, 10)
```

```
In [4]: test.shape
```

```
Out[4]: (58921, 10)
```

Train Data Analysis

```
In [5]: train.columns
```

```
Out[5]: Index(['ACTION', 'RESOURCE', 'MGR_ID', 'ROLE_ROLLUP_1', 'ROLE_ROLLUP_2',  
              'ROLE_DEPTNAME', 'ROLE_TITLE', 'ROLE_FAMILY_DESC', 'ROLE_FAMILY',  
              'ROLE_CODE'],  
             dtype='object')
```

```
In [6]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 32769 entries, 0 to 32768  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   ACTION                32769 non-null  int64  
1   RESOURCE              32769 non-null  int64  
2   MGR_ID                32769 non-null  int64  
3   ROLE_ROLLUP_1         32769 non-null  int64  
4   ROLE_ROLLUP_2         32769 non-null  int64  
5   ROLE_DEPTNAME         32769 non-null  int64  
6   ROLE_TITLE            32769 non-null  int64  
7   ROLE_FAMILY_DESC      32769 non-null  int64  
8   ROLE_FAMILY           32769 non-null  int64  
9   ROLE_CODE             32769 non-null  int64  
dtypes: int64(10)  
memory usage: 2.5 MB
```

```
In [7]: train.head()
```

```
Out[7]:
```

ACTION	RESOURCE	MGR_ID	ROLE_ROLLUP_1	ROLE_ROLLUP_2	ROLE_DEPTNAME	ROLE_FAMILY	ROLE_FAMILY_DESC	ROLE_TITLE	ROLE_CODE
--------	----------	--------	---------------	---------------	---------------	-------------	------------------	------------	-----------

	ACTION	RESOURCE	MGR_ID	ROLE_ROLLUP_1	ROLE_ROLLUP_2	ROLE_DEPTNAME	ROLE_D
0	1	39353	85475	117961	118300	123472	
1	1	17183	1540	117961	118343	123125	
2	1	36724	14457	118219	118220	117884	
3	1	36135	5396	117961	118343	119993	
4	1	42680	5905	117929	117930	119569	

In [8]: `train.describe()`

Out[8]:

	ACTION	RESOURCE	MGR_ID	ROLE_ROLLUP_1	ROLE_ROLLUP_2	ROLE_D
count	32769.000000	32769.000000	32769.000000	32769.000000	32769.000000	32769.000000
mean	0.942110	42923.916171	25988.957979	116952.627788	118301.823156	118301.823156
std	0.233539	34173.892702	35928.031650	10875.563591	4551.588572	1891.588572
min	0.000000	0.000000	25.000000	4292.000000	23779.000000	4600.000000
25%	1.000000	20299.000000	4566.000000	117961.000000	118102.000000	118301.823156
50%	1.000000	35376.000000	13545.000000	117961.000000	118300.000000	118301.823156
75%	1.000000	74189.000000	42034.000000	117961.000000	118386.000000	120500.000000
max	1.000000	312153.000000	311696.000000	311178.000000	286791.000000	286791.000000

In [9]: `# unique values
for i in train:
 print(i, len(train[i].unique()))`

ACTION 2
RESOURCE 7518
MGR_ID 4243
ROLE_ROLLUP_1 128
ROLE_ROLLUP_2 177
ROLE_DEPTNAME 449

```
ROLE_TITLE 343
ROLE_FAMILY_DESC 2358
ROLE_FAMILY 67
ROLE_CODE 343
```

ROLE_TITLE and ROLE_CODE columns has same no. of entries, In other words we can say both columns are same. ACTION is our class label.

```
In [10]: train.isna().sum()
```

```
Out[10]: ACTION          0
RESOURCE          0
MGR_ID            0
ROLE_ROLLUP_1     0
ROLE_ROLLUP_2     0
ROLE_DEPTNAME     0
ROLE_TITLE        0
ROLE_FAMILY_DESC  0
ROLE_FAMILY       0
ROLE_CODE         0
dtype: int64
```

```
In [11]: train.duplicated().sum()
```

```
Out[11]: 0
```

There is no duplicated and missing values in the train dataset

Test Data Analysis

```
In [12]: test.columns
```

```
Out[12]: Index(['id', 'RESOURCE', 'MGR_ID', 'ROLE_ROLLUP_1', 'ROLE_ROLLUP_2',
               'ROLE_DEPTNAME', 'ROLE_TITLE', 'ROLE_FAMILY_DESC', 'ROLE_FAMIL
               Y',
```

```
'ROLE_CODE'],  
dtype='object')
```

In [13]: `test.info()`

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 58921 entries, 0 to 58920  
Data columns (total 10 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   id                    58921 non-null  int64  
1   RESOURCE              58921 non-null  int64  
2   MGR_ID                58921 non-null  int64  
3   ROLE_ROLLUP_1         58921 non-null  int64  
4   ROLE_ROLLUP_2         58921 non-null  int64  
5   ROLE_DEPTNAME         58921 non-null  int64  
6   ROLE_TITLE            58921 non-null  int64  
7   ROLE_FAMILY_DESC      58921 non-null  int64  
8   ROLE_FAMILY           58921 non-null  int64  
9   ROLE_CODE             58921 non-null  int64  
dtypes: int64(10)  
memory usage: 4.5 MB
```

In [14]: `test.head()`

Out[14]:

	id	RESOURCE	MGR_ID	ROLE_ROLLUP_1	ROLE_ROLLUP_2	ROLE_DEPTNAME	ROLE_TITLE
0	1	78766	72734	118079	118080	117878	117879
1	2	40644	4378	117961	118327	118507	118863
2	3	75443	2395	117961	118300	119488	118172
3	4	43219	19986	117961	118225	118403	120773
4	5	42093	50015	117961	118343	119598	118422

In [15]: `test.describe()`

Out[15]:

	id	RESOURCE	MGR_ID	ROLE_ROLLUP_1	ROLE_ROLLUP_2	ROLE_D
count	58921.000000	58921.000000	58921.000000	58921.000000	58921.000000	58921.000000
mean	29461.000000	39383.739482	26691.645050	117028.638041	118316.334091	118316.334091
std	17009.171942	33717.397122	35110.244281	10805.446548	4284.678750	17009.171942
min	1.000000	0.000000	25.000000	4292.000000	23779.000000	4292.000000
25%	14731.000000	18418.000000	4663.000000	117961.000000	118096.000000	118096.000000
50%	29461.000000	33248.000000	14789.000000	117961.000000	118300.000000	118300.000000
75%	44191.000000	45481.000000	46512.000000	117961.000000	118386.000000	120419.000000
max	58921.000000	312136.000000	311779.000000	311178.000000	194897.000000	277619.000000

In [16]:

```
# unique value
for i in test:
    print(i, len(test[i].unique()))
```

```
id 58921
RESOURCE 4971
MGR_ID 4689
ROLE_ROLLUP_1 126
ROLE_ROLLUP_2 177
ROLE_DEPTNAME 466
ROLE_TITLE 351
ROLE_FAMILY_DESC 2749
ROLE_FAMILY 68
ROLE_CODE 351
```

In [17]:

```
test.isna().sum()
```

Out[17]:

```
id 0
RESOURCE 0
MGR_ID 0
ROLE_ROLLUP_1 0
ROLE_ROLLUP_2 0
ROLE_DEPTNAME 0
```



```
ROLE_TITLE      0
ROLE_FAMILY_DESC 0
ROLE_FAMILY      0
ROLE_CODE        0
dtype: int64
```

```
In [18]: test.duplicated().sum()
```

```
Out[18]: 0
```

There is no duplicated and missing values in the test dataset

Analysing Individual Columns

ACTION

```
In [19]: train['ACTION'].value_counts()
```

```
Out[19]: 1    30872
         0     1897
         Name: ACTION, dtype: int64
```

```
In [20]: approved_actions = train[train.ACTION==1]
```

```
In [21]: rejected_actions = train[train.ACTION==0]
```

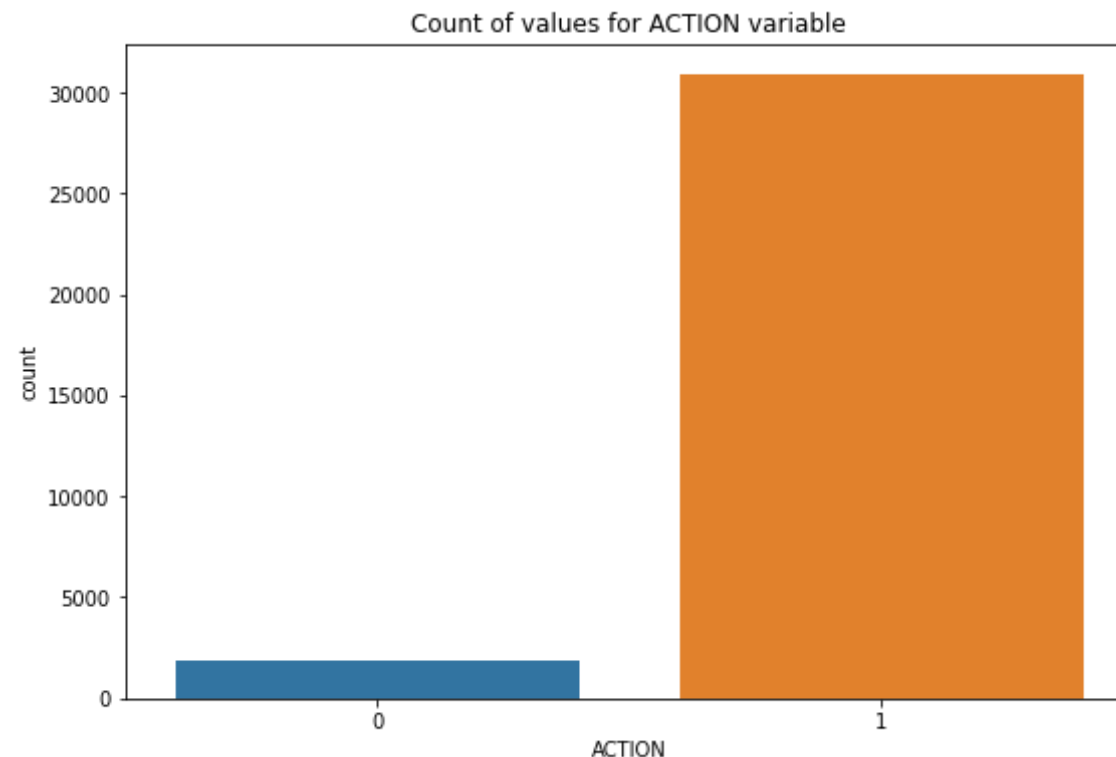
```
In [22]: approved_actions.shape
```

```
Out[22]: (30872, 10)
```

```
In [23]: rejected_actions.shape
```

```
Out[23]: (1897, 10)
```

```
In [24]: plt.figure(figsize=(9,6));  
sb.countplot(x='ACTION',data=train);  
plt.title('Count of values for ACTION variable');
```




As per the graph we have imbalanced data set, frequency of approved requests are much greater than rejected one. So we have to find out some ways to make this dataset balance.

RESOURCE

```
In [25]: plt.figure(figsize=(15,6));  
sb.kdeplot(approved_actions['RESOURCE'],label='Accepted',shade=True);  
sb.kdeplot(rejected_actions['RESOURCE'],label='Rejected',shade=True);  
plt.title('Distribution of RESOURCE variable');
```

```
plt.xlabel('RESOURCE');
plt.ylabel('Probability Density');
```



```
In [26]: # Top five approved requests
         approved_actions['RESOURCE'].value_counts()[:5]
```

```
Out[26]: 4675      836
         79092    468
         75078    405
         3853     398
         25993    390
         Name: RESOURCE, dtype: int64
```

```
In [27]: # Another Top five approved requests
         approved_actions['RESOURCE'].value_counts()[5:10]
```

```
Out[27]: 75834     294
         6977      283
         32270     279
         42085     237
         17308     236
         Name: RESOURCE, dtype: int64
```


```
In [28]: # Top five rejected requests
rejected_actions['RESOURCE'].value_counts()[:5]
```

```
Out[28]: 20897    42
         18072    29
         13878    22
         25993    19
         27416    19
         Name: RESOURCE, dtype: int64
```

```
In [29]: # Another Top five rejected requests
rejected_actions['RESOURCE'].value_counts()[5:10]
```

```
Out[29]: 7543     17
         79092    16
         32270    16
         6977     16
         32642    13
         Name: RESOURCE, dtype: int64
```

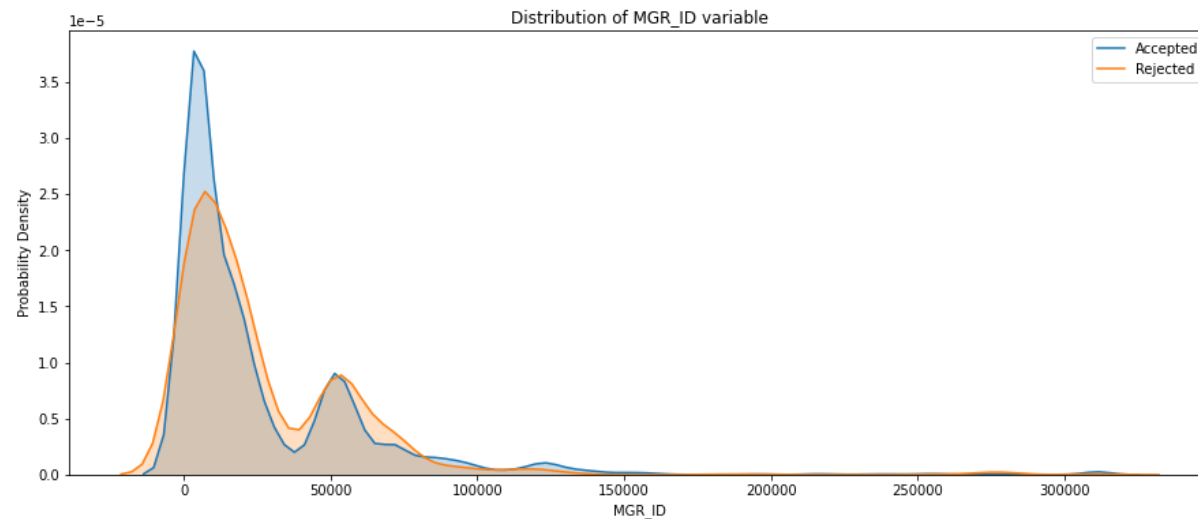
```
In [30]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['RESOURCE'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['RESOURCE'],label='Rejected',shade=True);
plt.title('Distribution of RESOURCE variable');
plt.xlim(0,100000)
plt.xlabel('RESOURCE');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 70K-90K Approved requests are higher than the rejected ones

MGR_ID

```
In [31]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['MGR_ID'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['MGR_ID'],label='Rejected',shade=True);
plt.title('Distribution of MGR_ID variable');
plt.xlabel('MGR_ID');
plt.ylabel('Probability Density');
```



```
In [32]: # Top 5 Approved Actions for attribute MGR_ID
approved_actions['MGR_ID'].value_counts()[:5]
```

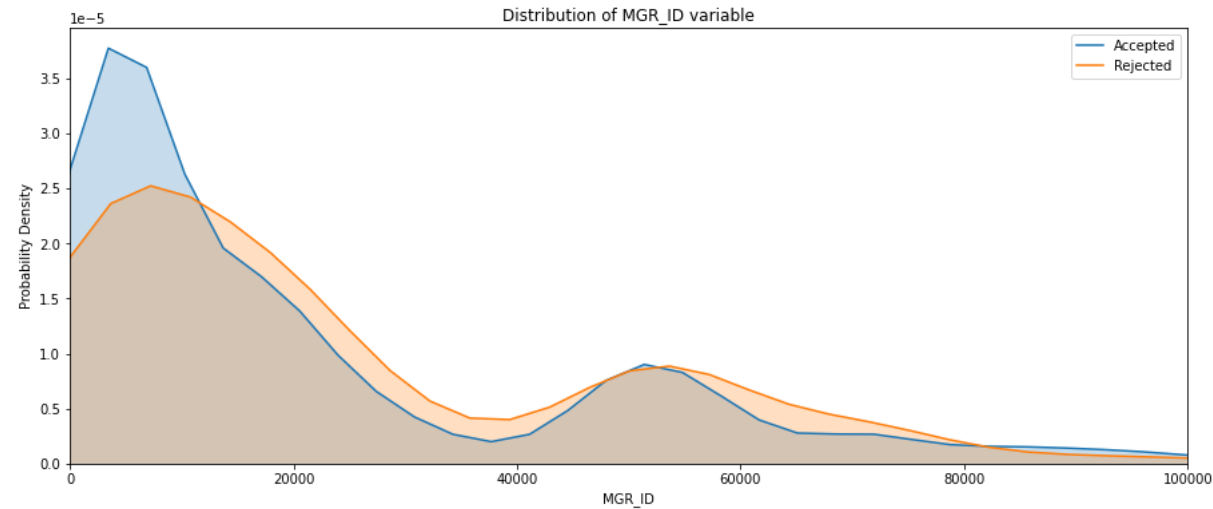
```
Out[32]: 770      147
2270      96
2594      71
2014      67
1350      67
Name: MGR_ID, dtype: int64
```

```
In [33]: # Top 5 Rejected Actions for attribute MGR_ID
rejected_actions['MGR_ID'].value_counts()[:5]
```

```
Out[33]: 54618     30
4084         17
46526        16
70062        16
4743         14
Name: MGR_ID, dtype: int64
```

```
In [34]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['MGR_ID'],label='Accepted',shade=True);
```

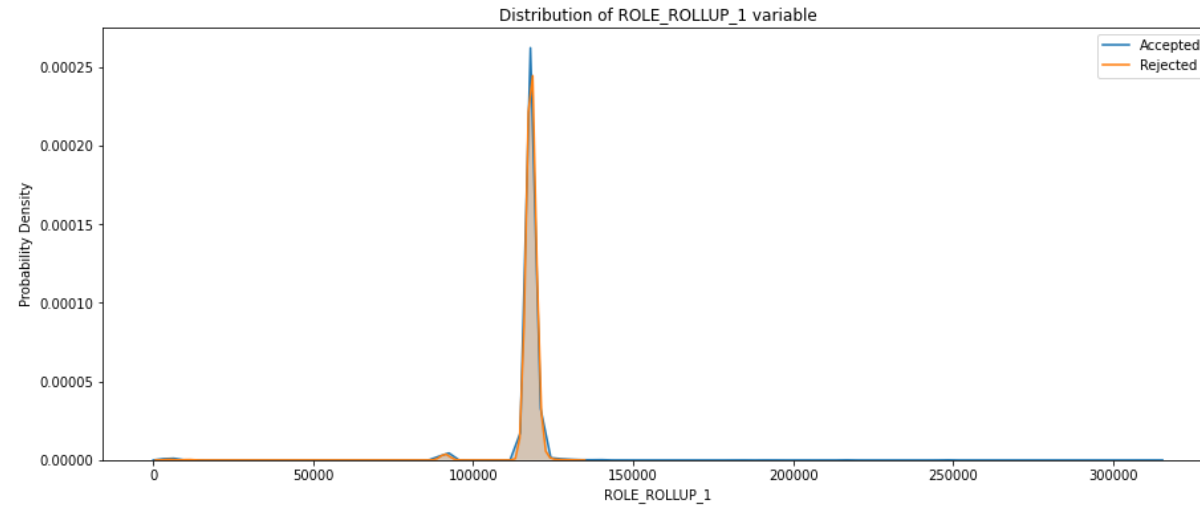
```
sb.kdeplot(rejected_actions['MGR_ID'],label='Rejected',shade=True);
plt.title('Distribution of MGR_ID variable');
plt.xlim(0,100000)
plt.xlabel('MGR_ID');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 0-20K Approved requests are higher than the rejected ones

ROLE_ROLLUP_1

```
In [35]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_ROLLUP_1'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_ROLLUP_1'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_ROLLUP_1 variable');
plt.xlabel('ROLE_ROLLUP_1');
plt.ylabel('Probability Density');
```



```
In [36]: # Top 5 Approved Actions for attribute ROLE_ROLLUP_1  
approved_actions['ROLE_ROLLUP_1'].value_counts()[:5]
```

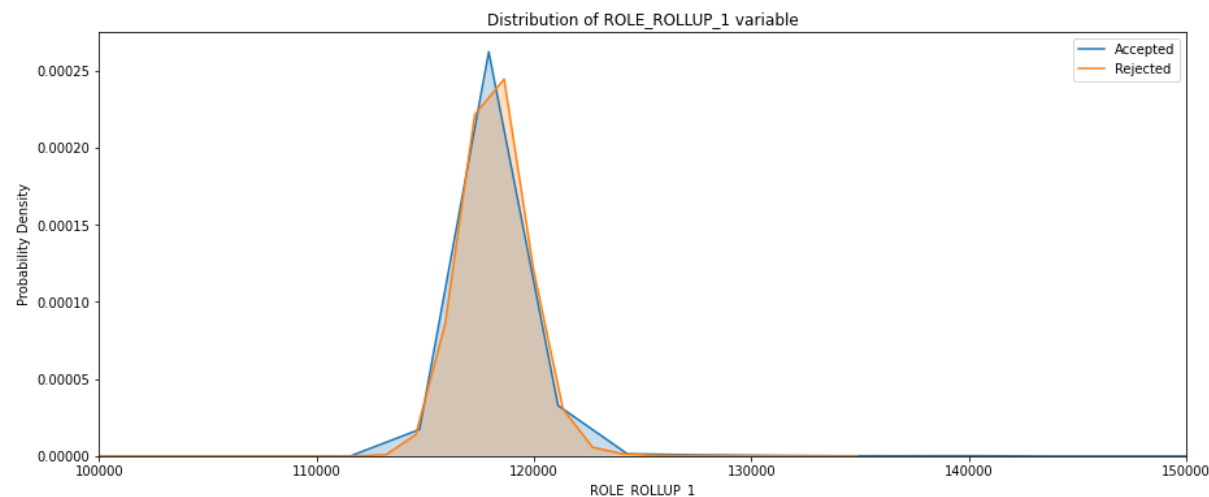
```
Out[36]: 117961    20320  
         117902     714  
         91261     695  
         118315    474  
         118212    385  
Name: ROLE_ROLLUP_1, dtype: int64
```

```
In [37]: # Top 5 Rejected Actions for attribute ROLE_ROLLUP_1  
rejected_actions['ROLE_ROLLUP_1'].value_counts()[:5]
```

```
Out[37]: 117961    1087  
         118256     73  
         119062     50  
         118290     44  
         118079     42  
Name: ROLE_ROLLUP_1, dtype: int64
```




```
In [38]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_ROLLUP_1'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_ROLLUP_1'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_ROLLUP_1 variable');
plt.xlim(100000,150000)
plt.xlabel('ROLE_ROLLUP_1');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that trends are almost similar

ROLE_ROLLUP_2

```
In [39]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_ROLLUP_2'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_ROLLUP_2'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_ROLLUP_2 variable');
plt.xlabel('ROLE_ROLLUP_2');
plt.ylabel('Probability Density');
```



```
In [40]: # Top 5 Approved Actions for attribute ROLE_ROLLUP_2
         approved_actions['ROLE_ROLLUP_2'].value_counts()[:5]
```

```
Out[40]: 118300    4230
         118343    3823
         118327    2521
         118225    2438
         118386    1639
         Name: ROLE_ROLLUP_2, dtype: int64
```

```
In [41]: # Top 5 Rejected Actions for attribute ROLE_ROLLUP_2
         rejected_actions['ROLE_ROLLUP_2'].value_counts()[:5]
```


```
Out[41]: 118300     194
         118052     185
         118386     157
         118343     122
         118327     120
         Name: ROLE_ROLLUP_2, dtype: int64
```

```
In [42]: plt.figure(figsize=(15,6));
```

```

sb.kdeplot(approved_actions['ROLE_ROLLUP_2'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_ROLLUP_2'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_ROLLUP_2 variable');
plt.xlim(100000, 150000)
plt.xlabel('ROLE_ROLLUP_2');
plt.ylabel('Probability Density');

```



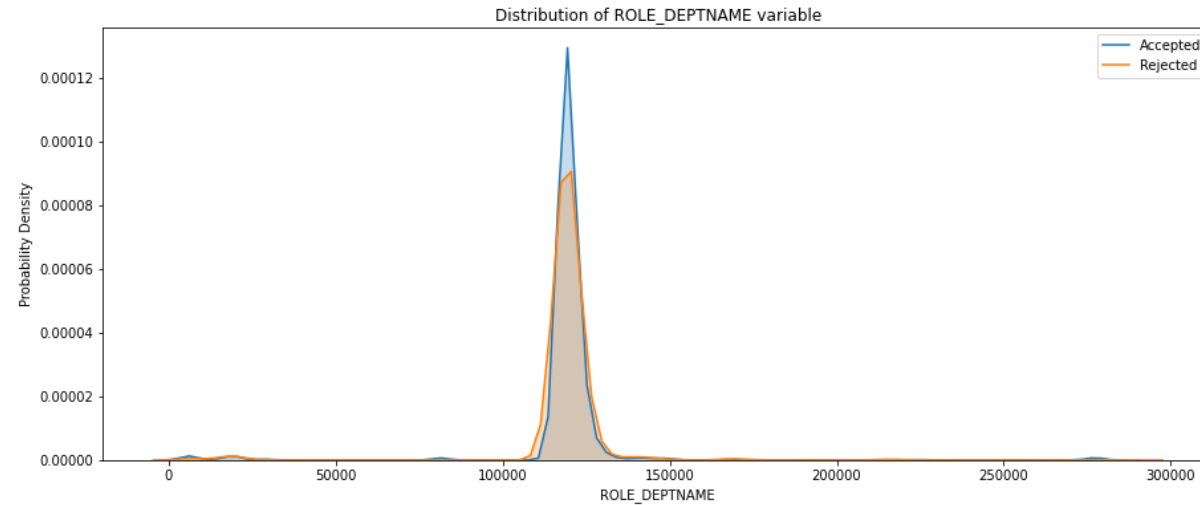
Looking at above KDE plot we can say that b/w 110K-120K Approved requests are higher than the rejected ones

ROLE_DEPTNAME

```

In [43]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_DEPTNAME'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_DEPTNAME'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_DEPTNAME variable');
plt.xlabel('ROLE_DEPTNAME');
plt.ylabel('Probability Density');

```



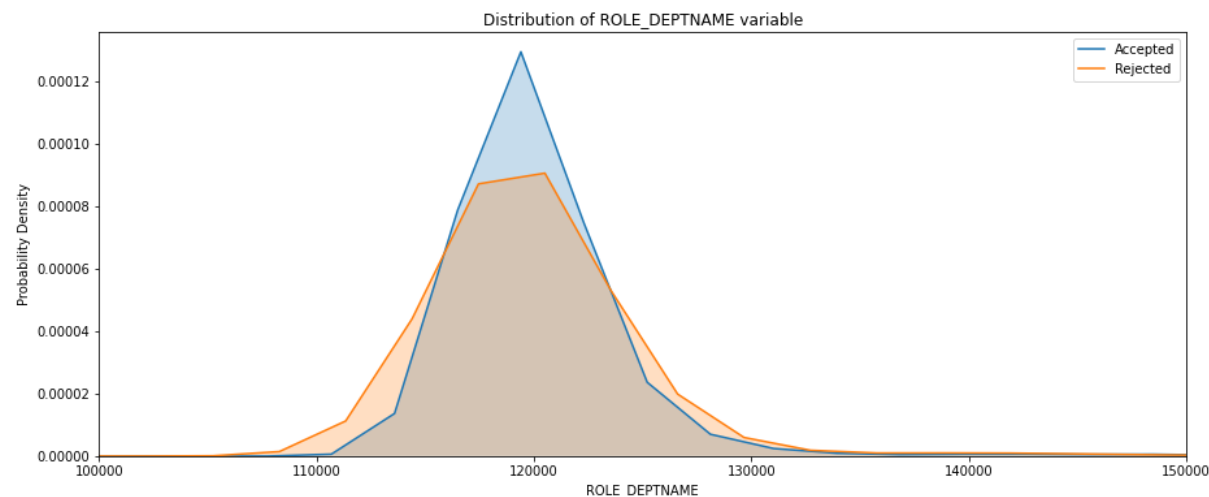
```
In [44]: # Top 5 Approved Actions for attribute ROLE_DEPTNAME
approved_actions['ROLE_DEPTNAME'].value_counts()[:5]
```

```
Out[44]: 117878    1064
         117941     700
         118514     589
         117945     570
         117920     541
         Name: ROLE_DEPTNAME, dtype: int64
```

```
In [45]: # Top 5 Rejected Actions for attribute ROLE_DEPTNAME
rejected_actions['ROLE_DEPTNAME'].value_counts()[:5]
```

```
Out[45]: 117945     89
         118992     77
         117878     71
         117941     63
         117920     56
         Name: ROLE_DEPTNAME, dtype: int64
```

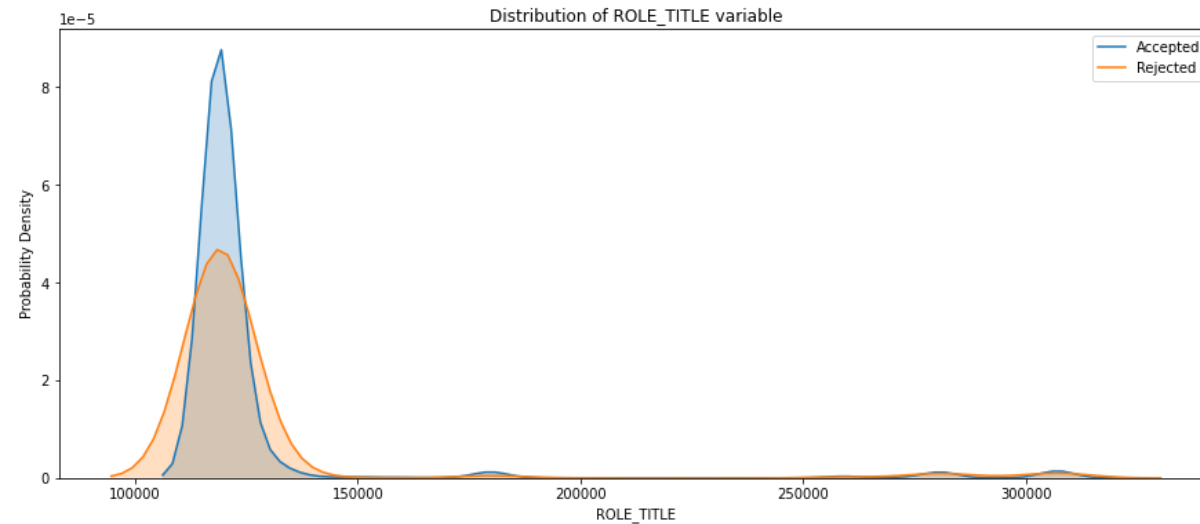
```
In [46]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_DEPTNAME'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_DEPTNAME'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_DEPTNAME variable');
plt.xlim(100000, 150000)
plt.xlabel('ROLE_DEPTNAME');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 110K-130K Approved requests are higher than the rejected ones

ROLE_TITLE

```
In [47]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_TITLE'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_TITLE'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_TITLE variable');
plt.xlabel('ROLE_TITLE');
plt.ylabel('Probability Density');
```



```
In [48]: # Top 5 Approved Actions for attribute ROLE_TITLE
approved_actions['ROLE_TITLE'].value_counts()[:5]
```


```
Out[48]: 118321    4279
         117905    3467
         118784    1647
         117879    1117
         118568     965
         Name: ROLE_TITLE, dtype: int64
```

```
In [49]: # Top 5 Rejected Actions for attribute ROLE_TITLE
rejected_actions['ROLE_TITLE'].value_counts()[:5]
```

```
Out[49]: 118321     370
         117879     139
         118784     125
         117905     116
         118568      78
         Name: ROLE_TITLE, dtype: int64
```

```
In [50]: plt.figure(figsize=(15,6));
         sb.kdeplot(approved_actions['ROLE_TITLE'],label='Accepted',shade=True);
```


```
sb.kdeplot(rejected_actions['ROLE_TITLE'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_TITLE variable');
plt.xlim(100000, 150000)
plt.xlabel('ROLE_TITLE');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 110K-130K Approved requests are higher than the rejected ones

ROLE_FAMILY_DESC

```
In [51]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_FAMILY_DESC'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_FAMILY_DESC'],label='Rejected',shade=True);
plt.title('Distribution of ROLE_FAMILY_DESC variable');
plt.xlabel('ROLE_FAMILY_DESC');
plt.ylabel('Probability Density');
```



```
In [52]: # Top 5 Approved Actions for attribute ROLE_FAMILY_DESC
         approved_actions['ROLE_FAMILY_DESC'].value_counts()[:5]
```

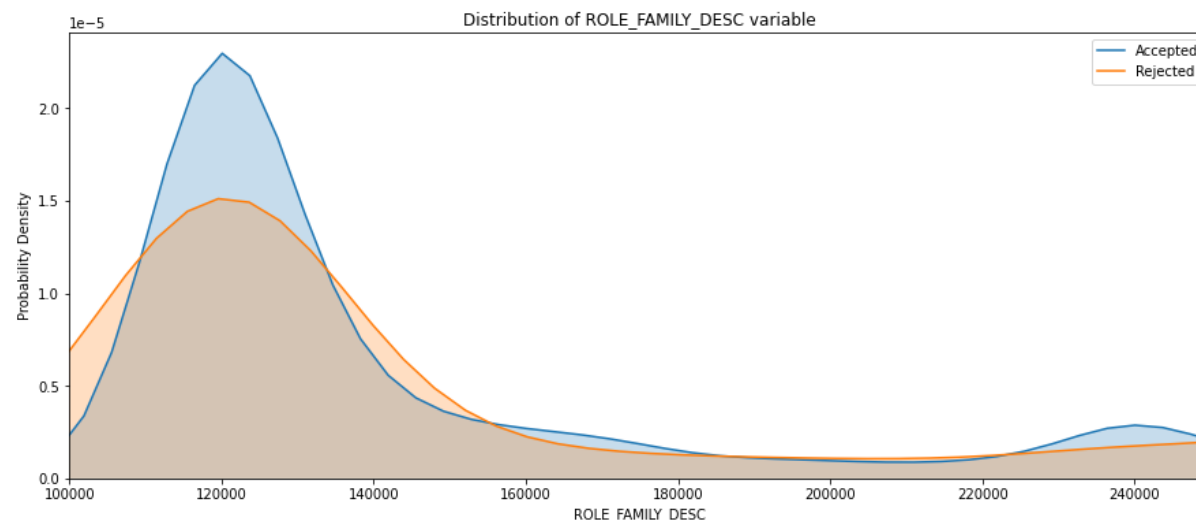
```
Out[52]: 117906      6437
         240983      1189
         117913       649
         279443       615
         117886       478
         Name: ROLE_FAMILY_DESC, dtype: int64
```

```
In [53]: # Top 5 Rejected Actions for attribute ROLE_FAMILY_DESC
         rejected_actions['ROLE_FAMILY_DESC'].value_counts()[:5]
```

```
Out[53]: 117906      459
         240983       55
         117886       52
         279443       50
         117897       39
         Name: ROLE_FAMILY_DESC, dtype: int64
```



```
In [54]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_FAMILY_DESC'],label='Accepted',shade=
True);
sb.kdeplot(rejected_actions['ROLE_FAMILY_DESC'],label='Rejected',shade=
True);
plt.title('Distribution of ROLE_FAMILY_DESC variable');
plt.xlim(100000, 250000)
plt.xlabel('ROLE_FAMILY_DESC');
plt.ylabel('Probability Density');
```




Looking at above KDE plot we can say that b/w 100K-140K Approved requests are higher than the rejected ones

ROLE_FAMILY

```
In [55]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_FAMILY'],label='Accepted',shade=True
);
sb.kdeplot(rejected_actions['ROLE_FAMILY'],label='Rejected',shade=True
);
plt.title('Distribution of ROLE_FAMILY variable');
```

```
plt.xlabel('ROLE_FAMILY');
plt.ylabel('Probability Density');
```



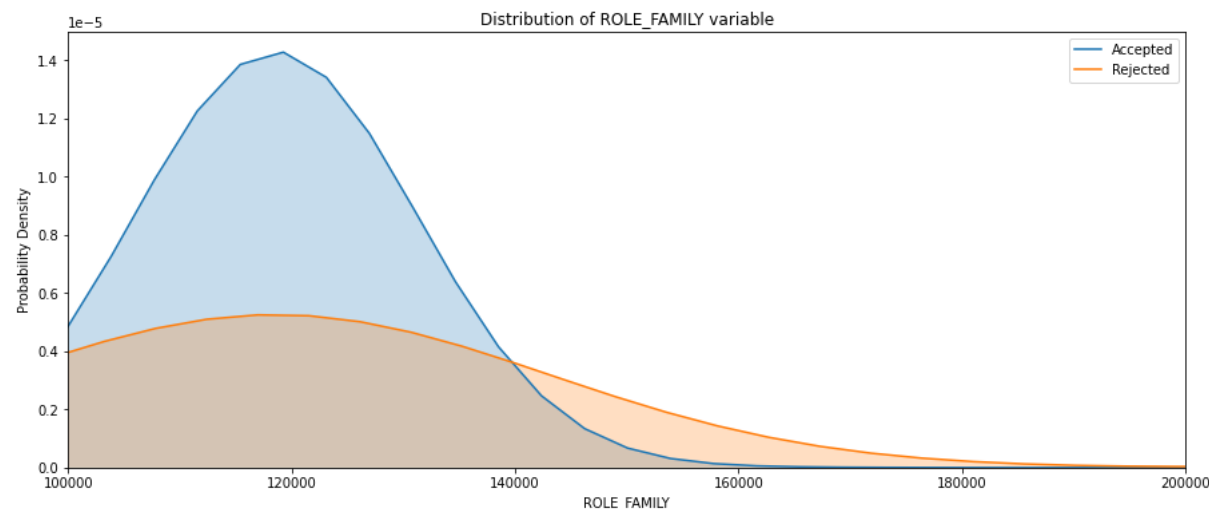
```
In [56]: # Top 5 Approved Actions for attribute ROLE_FAMILY
approved_actions['ROLE_FAMILY'].value_counts()[:5]
```

```
Out[56]: 290919    10347
118424      2616
19721       2393
117887      2302
118398      1232
Name: ROLE_FAMILY, dtype: int64
```

```
In [57]: # Top 5 Rejected Actions for attribute ROLE_FAMILY
rejected_actions['ROLE_FAMILY'].value_counts()[:5]
```


```
Out[57]: 290919      633
19721         243
292795       181
117887        98
118424        74
Name: ROLE_FAMILY, dtype: int64
```

```
In [58]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_FAMILY'],label='Accepted',shade=True
);
sb.kdeplot(rejected_actions['ROLE_FAMILY'],label='Rejected',shade=True
);
plt.title('Distribution of ROLE_FAMILY variable');
plt.xlim(100000, 200000)
plt.xlabel('ROLE_FAMILY');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 100K-140K Approved requests are higher than the rejected ones


```
In [59]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_FAMILY'],label='Accepted',shade=True
);
sb.kdeplot(rejected_actions['ROLE_FAMILY'],label='Rejected',shade=True
);
plt.title('Distribution of ROLE_FAMILY variable');
plt.xlim(200000, 300000)
plt.xlabel('ROLE_FAMILY');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w 260K-300K Approved requests are higher than the rejected ones

ROLE_CODE

```
In [60]: plt.figure(figsize=(15,6));  
sb.kdeplot(approved_actions['ROLE_CODE'],label='Accepted',shade=True);  
sb.kdeplot(rejected_actions['ROLE_CODE'],label='Rejected',shade=True);  
plt.title('Distribution of ROLE_CODE variable');  
plt.xlabel('ROLE_CODE');  
plt.ylabel('Probability Density');
```



```
In [61]: # Top 5 Approved Actions for attribute ROLE_CODE
approved_actions['ROLE_CODE'].value_counts()[:5]
```


```
Out[61]: 118322    4279
         117908    3467
         118786    1647
         117880    1117
         118570     965
         Name: ROLE_CODE, dtype: int64
```

```
In [62]: # Top 5 Rejected Actions for attribute ROLE_CODE
rejected_actions['ROLE_CODE'].value_counts()[:5]
```

```
Out[62]: 118322     370
         117880     139
         118786     125
         117908     116
         118570      78
         Name: ROLE_CODE, dtype: int64
```

```
In [63]: plt.figure(figsize=(15,6));
sb.kdeplot(approved_actions['ROLE_CODE'],label='Accepted',shade=True);
sb.kdeplot(rejected_actions['ROLE_CODE'],label='Rejected',shade=True);
```

```
plt.title('Distribution of ROLE_CODE variable');
plt.xlim(100000, 150000)
plt.xlabel('ROLE_CODE');
plt.ylabel('Probability Density');
```



Looking at above KDE plot we can say that b/w trends are almost similar for both classes

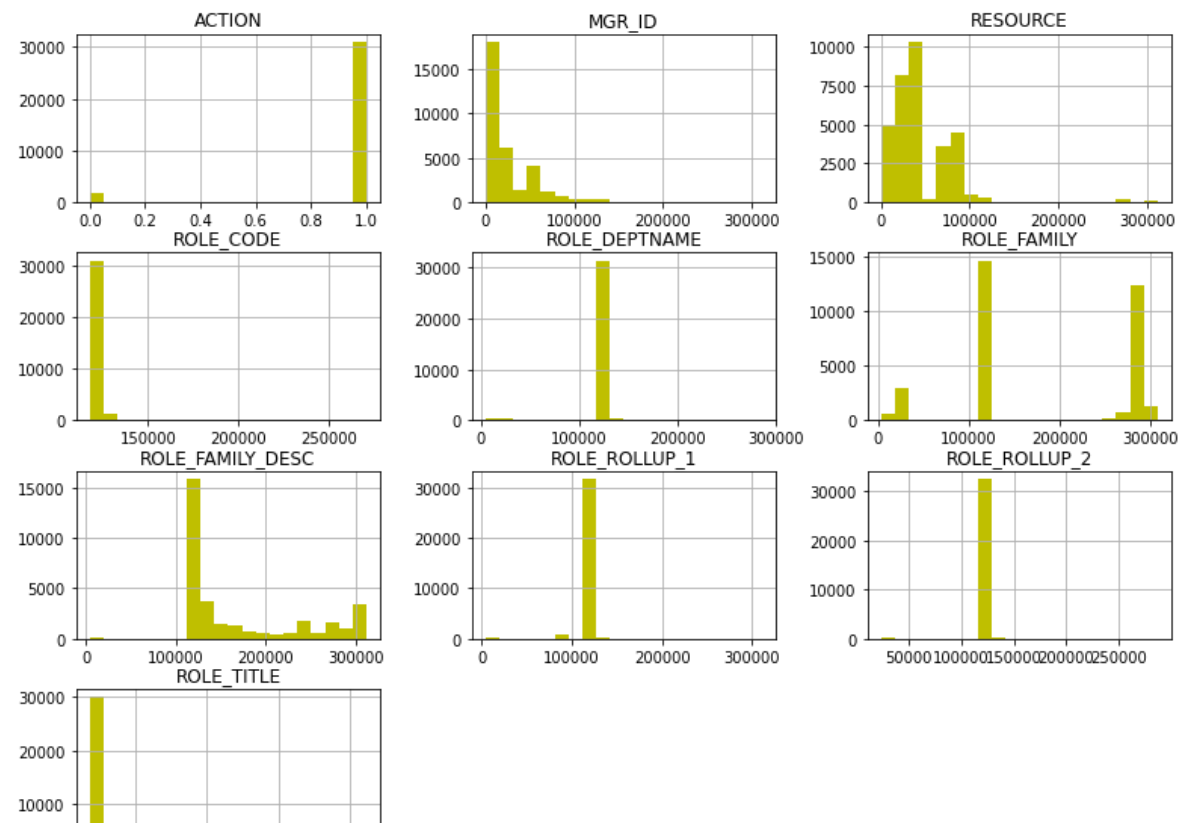
```
In [64]: train[['ACTION', 'RESOURCE', 'MGR_ID', 'ROLE_ROLLUP_1', 'ROLE_ROLLUP_2',
               'ROLE_DEPTNAME', 'ROLE_TITLE', 'ROLE_FAMILY_DESC', 'ROLE_FAMILY',
               'ROLE_CODE']].hist(figsize=(13,10),bins=20,color='Y')
```


```
Out[64]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215f11d
30>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x7f3216020a
90>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x7f3216021c
50>],
               [<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215fdb6
d8>,
               <matplotlib.axes._subplots.AxesSubplot object at 0x7f321a68d7
f0>],
               ...])
```

```


a0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f32160ad1
d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215ef21
98>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215ef25
50>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215ed21
d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215e769
e8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f3215df3b
e0>]],
dtype=object)

```





```
In [65]: plt.figure(figsize=(15,6));
sb.heatmap(train.corr(), annot=True, fmt='.2f');
plt.title('Heat Map between all features');
```



Observation

1. Almost all values are 0 expect correlation b/w (ROLE_FAMILY_DESC, ROLE_TITLE) and (ROLE_CODE, ROLE_TITLE)
2. Correlation b/w ROLE_FAMILY_DESC and ROLE_TITLE is 0.17
3. Correlation b/w ROLE_CODE and ROLE_TITLE is 0.16

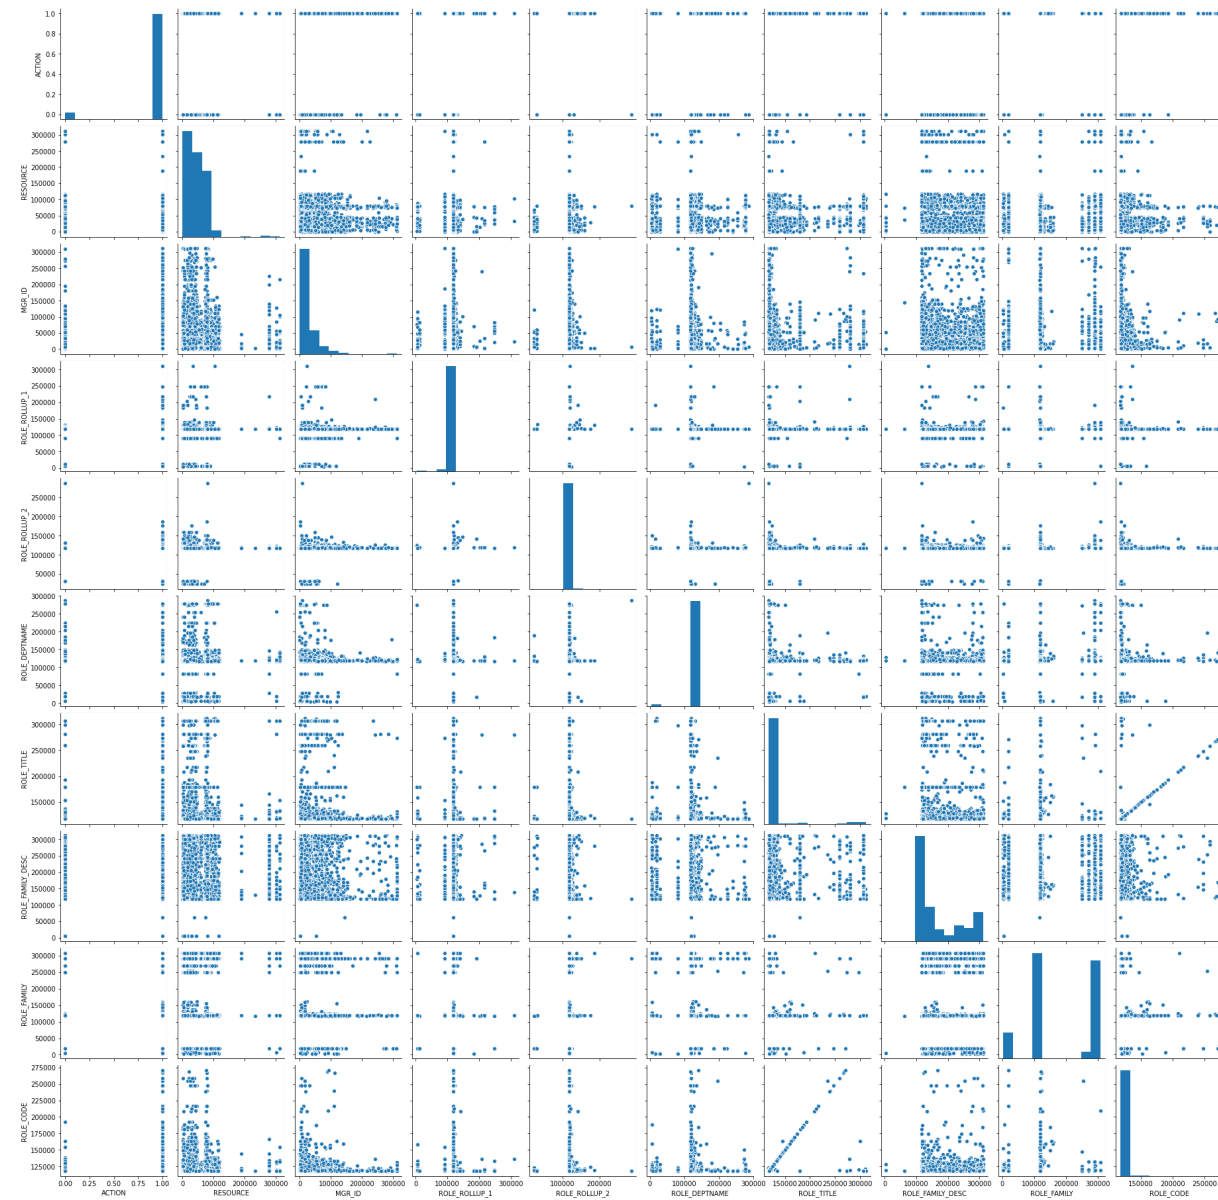
```
In [67]: plt.figure(figsize=(15,6))
```



```
sb.pairplot(train[['ACTION', 'RESOURCE', 'MGR_ID', 'ROLE_ROLLUP_1', 'ROLE_ROLLUP_2',  
                  'ROLE_DEPTNAME', 'ROLE_TITLE', 'ROLE_FAMILY_DESC', 'ROLE_FAMILY'  
,  
                  'ROLE_CODE']])
```

Out[67]: <seaborn.axisgrid.PairGrid at 0x7f320fd371d0>

<Figure size 1080x432 with 0 Axes>



Observation:

There is only relationship b/w ROLE_CODE and ROLE_TITLE

In []: