

```
In [1]: # import warnings
# warnings.filterwarnings('ignore')
```

```
In [2]: # import libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import sparse
%matplotlib inline

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import uniform
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from catboost import CatBoostClassifier
import pickle
```

Amazon Employee Access Challenge

```
In [3]: train = pd.read_csv('data/train.csv')
test = pd.read_csv('data/test.csv')
```

```
In [4]: train.shape
```

```
Out[4]: (32769, 10)
```

```
In [5]: test.shape
```

Out[5]: (58921, 10)

```
In [6]: y_train = train['ACTION']
```

```
In [7]: y_train.shape
```

Out[7]: (32769,)

```
In [8]: train_data = train.drop('ACTION', axis=1)
train_data.shape
```

Out[8]: (32769, 9)

```
In [9]: test_data = test.drop('id', axis=1)
test_data.shape
```

Out[9]: (58921, 9)

Common Variables

```
In [10]: # define variables
random_state = 42
cv = 5
scoring = 'roc_auc'
verbose=2
```

Common functions

```
In [11]: def save_submission(predictions, filename):
        '''
        Save predictions into csv file
        '''
        global test
        submission = pd.DataFrame()
```

```
submission["Id"] = test["id"]
submission["ACTION"] = predictions
filepath = "result/sampleSubmission_"+filename
submission.to_csv(filepath, index = False)
```

```
In [12]: def print_graph(results, param1, param2, xlabel, ylabel, title='Plot showing the ROC_AUC score for various hyper parameter values'):
        """
        Plot the graph
        """
        plt.plot(results[param1], results[param2]);
        plt.grid();
        plt.xlabel(xlabel);
        plt.ylabel(ylabel);
        plt.title(title);
```

```
In [13]: def get_rf_params():
        """
        Return dictionary of parameters for random forest
        """
        params = {
            'n_estimators': [10, 20, 50, 100, 200, 500, 700, 1000],
            'max_depth': [1, 2, 5, 10, 12, 15, 20, 25],
            'max_features': [1, 2, 3, 4, 5],
            'min_samples_split': [2, 5, 7, 10, 20]
        }

        return params
```

```
In [14]: def get_xgb_params():
        """
        Return dictionary of parameters for xgboost
        """
        params = {
            'n_estimators': [10, 20, 50, 100, 200, 500, 750, 1000],
            'learning_rate': uniform(0.01, 0.6),
            'subsample': uniform(),
            'max_depth': [3, 4, 5, 6, 7, 8, 9],
```

```
'colsample_bytree': uniform(),  
'min_child_weight': [1, 2, 3, 4]  
}  
  
return params
```

We will try following models

1. KNN
2. SVM
3. Logistic Regression
4. Random Forest
5. Xgboost

Build Models on the raw data

1.1 KNN with raw features

```
In [15]: parameters={'n_neighbors': np.arange(1, 100, 5)}  
clf = RandomizedSearchCV(KNeighborsClassifier(n_jobs=-1), parameters, ran  
dom_state=random_state, cv=cv, verbose=verbose, scoring=scoring, n_jobs=-1)  
best_model = clf.fit(train_data, y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work  
ers.  
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:      5.2s  
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:      8.8s finished
```

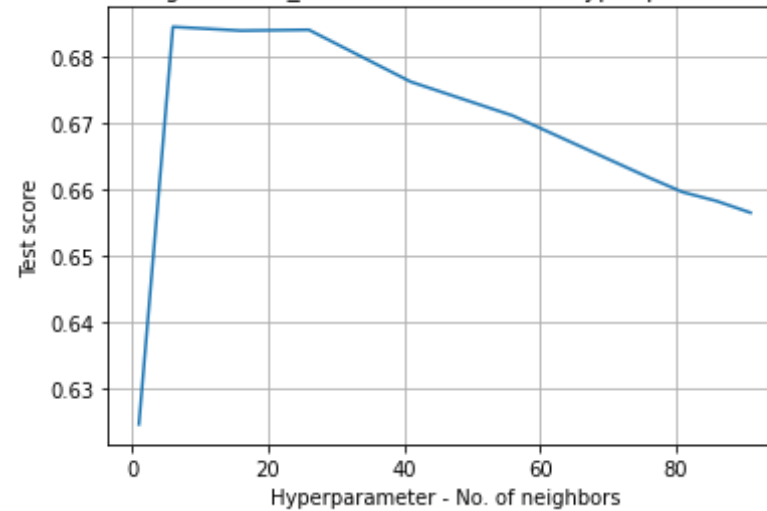
```
In [16]: results = pd.DataFrame.from_dict(best_model.cv_results_)  
results=results.sort_values('param_n_neighbors')  
results
```

Out[16]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	param
0	0.083934	0.007898	0.265385	0.036920	1	{'n_neighbors': 1}
3	0.240013	0.066697	0.500296	0.106391	6	{'n_neighbors': 6}
7	0.226630	0.185183	0.727930	0.213502	16	{'n_neighbors': 16}
5	0.243259	0.052800	0.871988	0.100550	26	{'n_neighbors': 26}
4	0.183503	0.043958	0.783068	0.125342	41	{'n_neighbors': 41}
6	0.228750	0.048025	1.059379	0.235010	56	{'n_neighbors': 56}
2	0.311753	0.040799	1.216632	0.265773	76	{'n_neighbors': 76}
9	0.270957	0.199804	0.948423	0.458374	81	{'n_neighbors': 81}
1	0.168152	0.078784	1.293272	0.219475	86	{'n_neighbors': 86}
8	0.108329	0.024517	1.590826	0.059955	91	{'n_neighbors': 91}

```
In [17]: print_graph(results, 'param_n_neighbors', 'mean_test_score', 'Hyperparameter - No. of neighbors', 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [18]: best_c=best_model.best_params_['n_neighbors']  
best_c
```

```
Out[18]: 6
```

```
In [19]: model = KNeighborsClassifier(n_neighbors=best_c,n_jobs=-1)  
model.fit(train_data,y_train)
```

```
Out[19]: KNeighborsClassifier(n_jobs=-1, n_neighbors=6)
```

```
In [20]: predictions = model.predict_proba(test_data)[: ,1]  
save_submission(predictions, "knn_raw.csv")
```

[sampleSubmission_knn_raw.csv](#)

20 hours ago by [Mayank Gupta](#)

I implemented KNN using only raw features.

0.67224

0.68148



1.2 SVM with raw feature

```
In [21]: C_val = uniform(loc=0, scale=4)
model= LinearSVC(verbose=verbose,random_state=random_state,class_weight
='balanced',max_iter=2000)
parameters={'C':C_val}
clf = RandomizedSearchCV(model,parameters,random_state=random_state,cv=
cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_data,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 2.0min finished
```

```
[LibLinear]
```

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n
umber of iterations.
"the number of iterations.", ConvergenceWarning)
```

```
In [22]: best_c=best_model.best_params_['C']
best_c
```

```
Out[22]: 1.49816047538945
```

```
In [23]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```


```
Out[23]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
6	18.684481	0.499781	0.009284	0.000217	0.232334	{'C': 0.23233444867279784}
5	19.870518	0.305372	0.009752	0.000904	0.623978	{'C': 0.6239780813448106}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
4	19.918902	0.358008	0.011202	0.001661	0.624075	{'C': 0.6240745617697461}
0	19.442288	0.216182	0.009220	0.000190	1.49816	{'C': 1.49816047538945}
3	19.357509	0.560751	0.009917	0.000562	2.39463	{'C': 2.3946339367881464}
8	18.831271	0.356759	0.008356	0.001407	2.40446	{'C': 2.404460046972835}
9	14.130057	5.223100	0.006211	0.001171	2.83229	{'C': 2.832290311184182}
2	18.946967	0.453250	0.009577	0.000543	2.92798	{'C': 2.9279757672456204}
7	18.603018	0.394303	0.009629	0.000678	3.4647	{'C': 3.4647045830997407}
1	19.380741	0.234885	0.009244	0.000622	3.80286	{'C': 3.8028572256396647}

```
In [24]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C',
, 'Test score')
```


Plot showing the ROC_AUC score for various hyper parameter values



In [25]: [#https://stackoverflow.com/questions/26478000/converting-linearsvcs-decision-function-to-probabilities-scikit-learn-python](https://stackoverflow.com/questions/26478000/converting-linearsvcs-decision-function-to-probabilities-scikit-learn-python)
model = LinearSVC(C=best_c, verbose=verbose, random_state=random_state, class_weight='balanced', max_iter=2000)
model = CalibratedClassifierCV(model)
model.fit(train_data, y_train)

[LibLinear]

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p  
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n  
umber of iterations.  
"the number of iterations.", ConvergenceWarning)
```

[LibLinear]

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p  
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n  
umber of iterations.  
"the number of iterations.", ConvergenceWarning)
```

[LibLinear]

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n
umber of iterations.
    "the number of iterations.", ConvergenceWarning)
```

```
[LibLinear]
```

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n
umber of iterations.
    "the number of iterations.", ConvergenceWarning)
```

```
[LibLinear]
```

```
/home/auw-mayank/.local/lib/python3.6/site-packages/sklearn/svm/_base.p
y:977: ConvergenceWarning: Liblinear failed to converge, increase the n
umber of iterations.
    "the number of iterations.", ConvergenceWarning)
```

```
Out[25]: CalibratedClassifierCV(base_estimator=LinearSVC(C=1.49816047538945,
class_weight='balance
d',
max_iter=2000, random_s
tate=42,
verbose=2))
```

```
In [26]: predictions = model.predict_proba(test_data)[: ,1]
save_submission(predictions, 'svm_raw.csv')
```

[sampleSubmission_svm_raw.csv](#)

0.50286

0.51390

20 hours ago by [Mayank Gupta](#)

I implemented SVM using only raw features.

1.3 Logistic Regression with Raw Feature

```
In [27]: C_val = uniform(loc=0, scale=4)
lr= LogisticRegression(verbose=verbose, random_state=random_state, class_
weight='balanced', solver='lbfgs', max_iter=500, n_jobs=-1)
parameters={'C':C_val}
clf = RandomizedSearchCV(lr, parameters, random_state=random_state, cv=cv,
verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model = clf.fit(train_data, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    1.2s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed:    6.1s
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed:   14.3s
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   20.5s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed:    0.3s finished
```

```
In [28]: best_c=best_model.best_params_['C']
best_c
```

Out[28]: 1.49816047538945


```
In [29]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```

Out[29]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_C
72	0.305453	0.042840	0.007123	0.001411	0.0220885	0.02208846849440959
10	0.299544	0.043052	0.009865	0.001636	0.082338	0.0823379771832097
98	0.324737	0.029356	0.009632	0.001840	0.101677	0.1016765069763807


	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param
42	0.322684	0.047769	0.008097	0.001223	0.137554	0.1375540844608739
58	0.312479	0.040692	0.010287	0.003491	0.180909	0.1809091556421522
...
1	0.330116	0.064433	0.008579	0.001581	3.80286	3.802857225639664
34	0.307352	0.038496	0.009020	0.000806	3.86253	3.862528132298237
50	0.286139	0.050608	0.008554	0.001624	3.87834	3.878338511058234
11	0.311766	0.046068	0.009899	0.002444	3.87964	3.879639408647977
69	0.288068	0.053532	0.009172	0.001290	3.94755	3.94754774640206

100 rows × 14 columns



```
In [30]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C'
, 'Test score')
```

Plot showing the ROC AUC score for various hyper parameter values



```
In [31]: model = LogisticRegression(C=best_c,verbose=verbose,n_jobs=-1,random_state=random_state,class_weight='balanced',solver='lbfgs')
model.fit(train_data,y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.3s finished

```
Out[31]: LogisticRegression(C=1.49816047538945, class_weight='balanced', n_jobs=-1,
                             random_state=42, verbose=2)
```

```
In [32]: predictions = model.predict_proba(test_data)[: ,1]
save_submission(predictions, 'lr_raw.csv')
```

[sampleSubmission_lr_raw.csv](#)

0.53857

0.53034

20 hours ago by [Mayank Gupta](#)

I implemented Logistic Regression using only raw features.

1.4 Random Forest with Raw Feature

```
In [33]: rfc = RandomForestClassifier(random_state=random_state, class_weight='balanced', n_jobs=-1)
         clf = RandomizedSearchCV(rfc, get_rf_params(), random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
         best_model = clf.fit(train_data, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 21.1s
/home/auw-mayank/.local/lib/python3.6/site-packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.
  "timeout or by a memory leak.", UserWarning
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 4.6min
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed: 10.3min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 15.1min finished
```

```
In [34]: results = pd.DataFrame(best_model.cv_results_)
         results.sort_values('mean_test_score', ascending=False, inplace=True)
         param_keys = ['param_' + str(each) for each in get_rf_params().keys()]
         param_keys.append('mean_test_score')
         results[param_keys].head(10)
```

Out[34]:

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
78	700	25	2	7	
62	500	25	3	5	
79	500	25	1	10	
55	200	25	2	5	
22	200	25	4	10	

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
20	1000	25	3	2	
85	1000	20	3	7	
33	700	25	4	2	
84	1000	25	5	2	
27	50	25	2	10	

```
In [35]: n_estimators=clf.best_params_['n_estimators']
max_features=clf.best_params_['max_features']
max_depth=clf.best_params_['max_depth']
min_samples_split=clf.best_params_['min_samples_split']
n_estimators,max_features,max_depth,min_samples_split
```

```
Out[35]: (700, 2, 25, 7)
```


```
In [36]: model=RandomForestClassifier(n_estimators=n_estimators,max_depth=max_de
pth,max_features=max_features,
                                     min_samples_split=min_samples_split,
                                     random_state=random_state,class_weight='ba
lanced',n_jobs=-1)

model.fit(train_data,y_train)
```

```
Out[36]: RandomForestClassifier(class_weight='balanced', max_depth=25, max_featu
res=2,
                                min_samples_split=7, n_estimators=700, n_jobs=-
1,
                                random_state=42)
```

```
In [37]: features=train_data.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features,'value':importance})
features=features.sort_values('value',ascending=False)
```

```
sns.barplot('value', 'features', data=features);  
plt.title('Feature Importance');
```



Features Observations:

1. MGR_ID is the most important feature followed by RESOURCE and ROLE_DEPTNAME

```
In [38]: predictions = model.predict_proba(test_data)[: ,1]  
save_submission(predictions, 'rf_raw.csv')
```

[sampleSubmission_rf_raw.csv](#)

0.87269

0.87567

20 hours ago by [Mayank Gupta](#)

I implemented Random Forest using only raw features.

1.5 Xgboost with Raw Feature


```
In [39]: xgb = XGBClassifier()
         clf = RandomizedSearchCV(xgb, get_xgb_params(), random_state=random_state,
                                cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
         best_model=clf.fit(train_data, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    7.2s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed:    1.5min
[Parallel(n_jobs=-1)]: Done 349 tasks    | elapsed:    4.1min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:    6.9min finished
```

```
In [40]: results = pd.DataFrame(best_model.cv_results_)
         results.sort_values('mean_test_score', ascending=False, inplace=True)
         param_keys=['param_'+str(each) for each in get_xgb_params().keys()]
         param_keys.append('mean_test_score')
         results[param_keys].head(10)
```

Out[40]:

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
18	1000	0.048135	0.665922	9	
44	1000	0.060484	0.606429	6	
97	750	0.232385	0.907694	6	
96	500	0.0979629	0.98664	7	
62	500	0.0663892	0.328153	9	
49	500	0.160277	0.393098	8	
84	200	0.571989	0.967581	6	
53	200	0.540096	0.928319	6	
86	1000	0.475848	0.858413	9	
8	750	0.0686033	0.683264	6	

```
In [41]: colsample_bytree = clf.best_params_['colsample_bytree']
learning_rate=clf.best_params_['learning_rate']
max_depth=clf.best_params_['max_depth']
min_child_weight=clf.best_params_['min_child_weight']
n_estimators=clf.best_params_['n_estimators']
subsample=clf.best_params_['subsample']
colsample_bytree, learning_rate, max_depth, min_child_weight, n_estimators,
subsample
```

```
Out[41]: (0.3308980248526492, 0.04813501017161418, 9, 2, 1000, 0.6659223566174967)
```

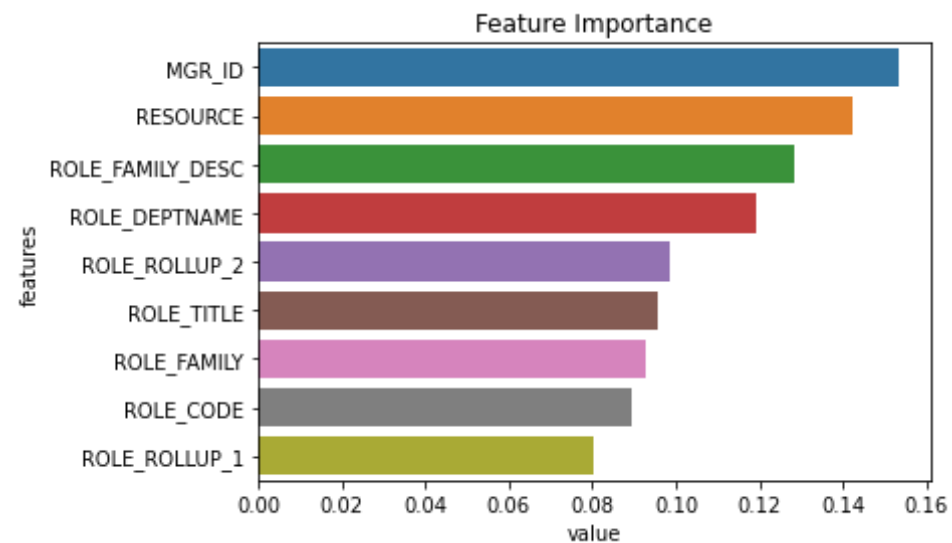
```
In [42]: model = XGBClassifier(colsample_bytree=colsample_bytree, learning_rate=learning_rate,
                               max_depth=max_depth,
                               min_child_weight=min_child_weight, n_estimators=n_estimators,
                               subsample=subsample, n_jobs=-1)

model.fit(train_data, y_train)
```

```
Out[42]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                       colsample_bynode=1, colsample_bytree=0.3308980248526492,
                       gamma=0,
                       gpu_id=-1, importance_type='gain', interaction_constraint
s='',
                       learning_rate=0.04813501017161418, max_delta_step=0, max_
depth=9,
                       min_child_weight=2, missing=nan, monotone_constraints
='()',
                       n_estimators=1000, n_jobs=-1, num_parallel_tree=1, random
_state=0,
                       reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                       subsample=0.6659223566174967, tree_method='exact',
                       validate_parameters=1, verbosity=None)
```

```
In [43]: features=train_data.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features, 'value':importance})
features=features.sort_values('value', ascending=False)
```

```
sns.barplot('value', 'features', data=features);  
plt.title('Feature Importance');
```



```
In [44]: predictions = model.predict_proba(test_data)[: ,1]  
save_submission(predictions, 'xgb_raw.csv')
```

[sampleSubmission_xgb_raw.csv](#)

0.86988

0.87909

20 hours ago by Mayank Gupta

I implemented Xgboost using only raw features.

Overview	Data	Notebooks	Discussion	Leaderboard	Rules	Team	My Submissions	Late Submission
Submission and Description						Private Score	Public Score	Use for Final Score
sampleSubmission_xgb_raw.csv 8 hours ago by Mayank Gupta I implemented Xgboost using only raw features.						0.86988	0.87909	<input type="checkbox"/>
sampleSubmission_rf_raw.csv 8 hours ago by Mayank Gupta I implemented Random Forest using only raw features.						0.87269	0.87567	<input type="checkbox"/>
sampleSubmission_lr_raw.csv 8 hours ago by Mayank Gupta I implemented Logistic Regression using only raw features.						0.53857	0.53034	<input type="checkbox"/>
sampleSubmission_svm_raw.csv 9 hours ago by Mayank Gupta I implemented SVM using only raw features.						0.50286	0.51390	<input type="checkbox"/>
sampleSubmission_knn_raw.csv 9 hours ago by Mayank Gupta I implemented KNN using only raw features.						0.67224	0.68148	<input type="checkbox"/>

```
In [45]: from prettytable import PrettyTable

x = PrettyTable(['Model', 'Feature', 'Private Score', 'Public Score'])
x.add_row(['KNN', 'Raw', 0.67224, 0.68148])
x.add_row(['SVM', 'Raw', 0.50286, 0.51390])
x.add_row(['Logistic Regression', 'Raw', 0.53857, 0.53034])
x.add_row(['Random Forest', 'Raw', 0.87269, 0.87567])
x.add_row(['Xgboost', 'Raw', 0.86988, 0.87909])

print(x)
```

Model	Feature	Private Score	Public Score
KNN	Raw	0.67224	0.68148
SVM	Raw	0.50286	0.5139
Logistic Regression	Raw	0.53857	0.53034
Random Forest	Raw	0.87269	0.87567

	Xgboost		Raw		0.86988		0.87909	
+-----+		+-----+		+-----+		+-----+		+-----+

Observations:

1. Xgboost perform best on the raw features
2. Random forest also perform good on raw features
3. Tree based models performs better than linear models for raw features

Build model on one hot encoded features

2.1 KNN with one hot encoded features

```
In [46]: train_ohe = sparse.load_npz('data/train_ohe.npz')
test_ohe = sparse.load_npz('data/test_ohe.npz')

train_ohe.shape, test_ohe.shape, y_train.shape
```

```
Out[46]: ((32769, 4500), (58921, 4500), (32769,))
```

```
In [47]: parameters={'n_neighbors':np.arange(1,100, 5)}
clf = RandomizedSearchCV(KNeighborsClassifier(n_jobs=-1),parameters,ran
dom_state=random_state,cv=cv,verbose=verbose,scoring=scoring,n_jobs=4)
best_model = clf.fit(train_ohe,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent worke
rs.
[Parallel(n_jobs=4)]: Done 33 tasks      | elapsed: 2.4min
[Parallel(n_jobs=4)]: Done 50 out of 50 | elapsed: 3.2min finished
```


```
In [48]: results = pd.DataFrame.from_dict(best_model.cv_results_)
```

```
results=results.sort_values('param_n_neighbors')
results
```

Out[48]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	param
0	0.008579	0.000699	12.889389	1.630046	1	{'n_neighbors': 1}
3	0.218653	0.141987	13.355838	1.399583	6	{'n_neighbors': 6}
7	0.014216	0.010986	11.480899	0.587346	16	{'n_neighbors': 16}
5	0.007408	0.000506	11.399296	0.072001	26	{'n_neighbors': 26}
4	0.029331	0.037009	11.730860	0.352833	41	{'n_neighbors': 41}
6	0.017152	0.013375	12.010613	0.940570	56	{'n_neighbors': 56}
2	0.325559	0.159848	31.339497	13.256170	76	{'n_neighbors': 76}
9	0.013539	0.006507	9.540637	2.054831	81	{'n_neighbors': 81}
1	0.103347	0.155776	24.211115	9.506957	86	{'n_neighbors': 86}
8	0.016657	0.008175	11.831052	0.749635	91	{'n_neighbors': 91}

In [49]: `print_graph(results, 'param_n_neighbors', 'mean_test_score', 'Hyperparameter - No. of neighbors', 'Test score')`



```
In [50]: best_c=best_model.best_params_['n_neighbors']  
best_c
```

```
Out[50]: 16
```

```
In [51]: model = KNeighborsClassifier(n_neighbors=best_c,n_jobs=-1)  
model.fit(train_ohe,y_train)
```

```
Out[51]: KNeighborsClassifier(n_jobs=-1, n_neighbors=16)
```

```
In [52]: predictions = model.predict_proba(test_ohe)[: ,1]  
save_submission(predictions, "knn_ohe.csv")
```

[sampleSubmission_knn_ohe.csv](#)

12 hours ago by Mayank Gupta

[add submission details](#)

0.81657

0.81723



2.2 SVM with one hot encoded features

```
In [53]: C_val = uniform(loc=0, scale=4)
model= LinearSVC(verbose=verbose,random_state=random_state,class_weight
='balanced',max_iter=2000)
parameters={'C':C_val}
clf = RandomizedSearchCV(model,parameters,random_state=random_state,cv=
cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_ohe,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    31.0s
```

```
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    41.7s finished
```

```
[LibLinear]
```

```
In [54]: best_c=best_model.best_params_['C']
best_c
```

```
Out[54]: 0.23233444867279784
```


```
In [55]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```

```
Out[55]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
6	1.730894	0.069332	0.005605	0.000029	0.232334	{'C': 0.23233444867279784}
5	3.315943	0.201073	0.005879	0.000622	0.623978	{'C': 0.6239780813448106}
4	3.583880	0.365580	0.005711	0.000297	0.624075	{'C': 0.6240745617697461}
0	5.492834	0.287928	0.006418	0.000139	1.49816	{'C': 1.49816047538945}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
3	4.961631	0.298074	0.005662	0.000146	2.39463	{'C': 2.3946339367881464}
8	4.833444	0.644881	0.004830	0.000888	2.40446	{'C': 2.404460046972835}
9	3.986315	0.360801	0.003652	0.000362	2.83229	{'C': 2.832290311184182}
2	4.873211	0.174700	0.006166	0.000985	2.92798	{'C': 2.9279757672456204}
7	4.817450	0.265487	0.005635	0.000099	3.4647	{'C': 3.4647045830997407}
1	4.685154	0.254005	0.006274	0.000241	3.80286	{'C': 3.8028572256396647}

```
In [56]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C', 'Test score')
```



```
In [57]: #https://stackoverflow.com/questions/26478000/converting-linearsvcs-dec
```

ision-function-to-probabilities-sci-kit-learn-python

```
model = LinearSVC(C=best_c, verbose=verbose, random_state=random_state, class_weight='balanced', max_iter=2000)
model = CalibratedClassifierCV(model)
model.fit(train_ohe, y_train)
```

```
[LibLinear][LibLinear][LibLinear][LibLinear][LibLinear]
```

```
Out[57]: CalibratedClassifierCV(base_estimator=LinearSVC(C=0.23233444867279784,
class_weight='balanced',
max_iter=2000, random_state=42,
verbose=2))
```

```
In [58]: predictions = model.predict_proba(test_ohe)[: ,1]
save_submission(predictions, 'svm_ohe.csv')
```

[sampleSubmission_svm_ohe.csv](#)

0.87249

0.87955



12 hours ago by [Mayank Gupta](#)

I tried models using one hot encoding of categorical variables.

2.3 Logistic Regression with one hot encoded features

```
In [59]: C_val = uniform(loc=0, scale=4)
lr= LogisticRegression(verbose=verbose, random_state=random_state, class_weight='balanced', solver='lbfgs', max_iter=500, n_jobs=-1)
parameters={'C':C_val}
clf = RandomizedSearchCV(lr, parameters, random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model = clf.fit(train_ohe, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work

```

ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    7.2s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed:   32.8s
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed:   1.4min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   2.0min finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed:    0.8s finished

```

```

In [60]: best_c=best_model.best_params_['C']
best_c

```

```

Out[60]: 0.6820964947491661

```

```

In [61]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results

```

```

Out[61]:

```


	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_
72	0.439680	0.041474	0.006678	0.000959	0.0220885	0.02208846849440959
10	0.650064	0.051074	0.006208	0.001374	0.082338	0.0823379771832097
98	0.776743	0.046280	0.006456	0.000961	0.101677	0.1016765069763807
42	0.886947	0.055574	0.008943	0.003785	0.137554	0.1375540844608739
58	0.928974	0.072238	0.006947	0.000977	0.180909	0.1809091556421522
...
1	2.534207	0.214963	0.006663	0.000782	3.80286	3.802857225639664
34	2.376340	0.051137	0.006265	0.000269	3.86253	3.862528132298237

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_C
50	2.498738	0.052591	0.007061	0.000964	3.87834	3.878338511058234
11	2.401271	0.107580	0.006628	0.001671	3.87964	3.879639408647977
69	2.674097	0.153710	0.007042	0.000979	3.94755	3.94754774640206

100 rows × 14 columns

```
In [62]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C',
, 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [63]: model = LogisticRegression(C=best_c,verbose=verbose,n_jobs=-1,random_state=random_state,class_weight='balanced',solver='lbfgs')
model.fit(train_ohe,y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.5s finished
```

```
Out[63]: LogisticRegression(C=0.6820964947491661, class_weight='balanced', n_jobs=-1,  
                             random_state=42, verbose=2)
```

```
In [64]: predictions = model.predict_proba(test_ohe)[: ,1]  
         save_submission(predictions, 'lr_ohe.csv')
```

[sampleSubmission_lr_ohe.csv](#)

12 hours ago by Mayank Gupta

[add submission details](#)

0.87436

0.88167



2.4 Random Forest with one hot encoded features

```
In [65]: rfc = RandomForestClassifier(random_state=random_state, class_weight='balanced', n_jobs=-1)  
         clf = RandomizedSearchCV(rfc, get_rf_params(), random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)  
         best_model = clf.fit(train_ohe, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks | elapsed: 9.8s
```

```
[Parallel(n_jobs=-1)]: Done 146 tasks | elapsed: 1.6min
```

```
[Parallel(n_jobs=-1)]: Done 349 tasks | elapsed: 3.6min
```

```
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 5.3min finished
```

```
In [66]: results = pd.DataFrame(best_model.cv_results_)  
         results.sort_values('mean_test_score', ascending=False, inplace=True)  
         param_keys = ['param_' + str(each) for each in get_rf_params().keys()]  
         param_keys.append('mean_test_score')  
         results[param_keys].head(10)
```

Out[66]:

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
78	700	25	2	7	
85	1000	20	3	7	
62	500	25	3	5	
6	500	20	2	5	
11	1000	15	3	7	
25	700	15	4	7	
19	700	15	4	5	
22	200	25	4	10	
79	500	25	1	10	
82	700	20	5	20	

In [67]:

```

n_estimators=clf.best_params_['n_estimators']
max_features=clf.best_params_['max_features']
max_depth=clf.best_params_['max_depth']
min_samples_split=clf.best_params_['min_samples_split']
n_estimators,max_features,max_depth,min_samples_split

```

Out[67]: (700, 2, 25, 7)

In [68]:

```

model=RandomForestClassifier(n_estimators=n_estimators,max_depth=max_de
pth,max_features=max_features,
                             min_samples_split=min_samples_split,
                             random_state=random_state,class_weight='ba
lanced',n_jobs=-1)

model.fit(train_ohe,y_train)

```

Out[68]: RandomForestClassifier(class_weight='balanced', max_depth=25, max_features=2, min_samples_split=7, n_estimators=700, n_jobs=-1, random_state=42)

```
In [69]: # features=train_ohe.columns
# importance=model.feature_importances_
# features=pd.DataFrame({'features':features,'value':importance})
# features=features.sort_values('value',ascending=False)
# sns.barplot('value','features',data=features);
# plt.title('Feature Importance');
```

```
In [70]: predictions = model.predict_proba(test_ohe)[:,:1]
save_submission(predictions, 'rf_ohe.csv')
```

[sampleSubmission_rf_ohe.csv](#)

0.84541

0.84997



12 hours ago by [Mayank Gupta](#)

[add submission details](#)

2.5 Xgboost with one hot encoded features

```
In [71]: xgb = XGBClassifier()
clf = RandomizedSearchCV(xgb,get_xgb_params(),random_state=random_state
,cv=cv,verbose=verbose,n_iter=100,scoring=scoring,n_jobs=-1)
best_model=clf.fit(train_ohe,y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    8.0s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed:   1.6min
[Parallel(n_jobs=-1)]: Done 349 tasks    | elapsed:   4.3min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   7.1min finished
```

```
In [72]: results = pd.DataFrame(best_model.cv_results_)
results.sort_values('mean_test_score',ascending=False,inplace=True)
param_keys=['param_'+str(each) for each in get_xgb_params().keys()]
```

```
param_keys.append('mean_test_score')
results[param_keys].head(10)
```

Out[72]:

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
97	750	0.232385	0.907694	6	
80	1000	0.385564	0.905351	3	
86	1000	0.475848	0.858413	9	
84	200	0.571989	0.967581	6	
14	200	0.374221	0.802197	7	
50	500	0.388683	0.645103	4	
53	200	0.540096	0.928319	6	
92	200	0.478778	0.49442	9	
96	500	0.0979629	0.98664	7	
22	1000	0.391846	0.695516	6	

```
In [73]: colsample_bytree = clf.best_params_['colsample_bytree']
learning_rate=clf.best_params_['learning_rate']
max_depth=clf.best_params_['max_depth']
min_child_weight=clf.best_params_['min_child_weight']
n_estimators=clf.best_params_['n_estimators']
subsample=clf.best_params_['subsample']
colsample_bytree, learning_rate, max_depth, min_child_weight, n_estimators,
subsample
```

Out[73]: (0.3742707957561203, 0.23238528824013455, 6, 1, 750, 0.9076937063485463)

```
In [74]: model = XGBClassifier(colsample_bytree=colsample_bytree, learning_rate=l
earning_rate, max_depth=max_depth,
min_child_weight=min_child_weight, n_estimators=n_e
stimators, subsample=subsample, n_jobs=-1)
```



```
model.fit(train_ohe,y_train)
```

```
Out[74]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=0.3742707957561203,
                        gamma=0, gpu_id=-1, importance_type='gain', interaction_constraint
                        s='', learning_rate=0.23238528824013455, max_delta_step=0, max_
                        depth=6, min_child_weight=1, missing=nan, monotone_constraints
                        ='()', n_estimators=750, n_jobs=-1, num_parallel_tree=1, random_
                        state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                        subsample=0.9076937063485463, tree_method='exact',
                        validate_parameters=1, verbosity=None)
```

```
In [75]: # features=train_ohe.columns
# importance=model.feature_importances_
# features=pd.DataFrame({'features':features,'value':importance})
# features=features.sort_values('value',ascending=False)
# sns.barplot('value','features',data=features);
# plt.title('Feature Importance');
```

```
In [76]: predictions = model.predict_proba(test_ohe)[:,:1]
save_submission(predictions, 'xgb_ohe.csv')
```

[sampleSubmission_xgb_ohe.csv](#)

12 hours ago by Mayank Gupta

[add submission details](#)

0.84717

0.85102



Overview	Data	Notebooks	Discussion	Leaderboard	Rules	Team	My Submissions	Late Submission
Submission and Description		Private Score		Public Score		Use for Final Score		
sampleSubmission_xgb_ohe.csv just now by Mayank Gupta add submission details		0.84717		0.85102		<input type="checkbox"/>		
sampleSubmission_rf_ohe.csv a minute ago by Mayank Gupta add submission details		0.84541		0.84997		<input type="checkbox"/>		
sampleSubmission_lr_ohe.csv 2 minutes ago by Mayank Gupta add submission details		0.87436		0.88167		<input type="checkbox"/>		
sampleSubmission_knn_ohe.csv 2 minutes ago by Mayank Gupta add submission details		0.81657		0.81723		<input type="checkbox"/>		
sampleSubmission_svm_ohe.csv 4 minutes ago by Mayank Gupta I tried models using one hot encoding of categorical variables.		0.87249		0.87955		<input type="checkbox"/>		

```
In [77]: from prettytable import PrettyTable

x = PrettyTable(['Model', 'Feature', 'Private Score', 'Public Score'])
x.add_row(['KNN', 'ohe', 0.81657, 0.81723])
x.add_row(['SVM', 'ohe', 0.87249, 0.87955])
x.add_row(['Logistic Regression', 'ohe', 0.87436, 0.88167])
x.add_row(['Random Forest', 'ohe', 0.84541, 0.84997])
x.add_row(['Xgboost', 'ohe', 0.84717, 0.85102])

print(x)
```

```
+-----+-----+-----+-----+
|      Model      | Feature | Private Score | Public Score |
+-----+-----+-----+-----+
|      KNN       | ohe    | 0.81657      | 0.81723      |
|      SVM       | ohe    | 0.87249      | 0.87955      |
| Logistic Regression | ohe    | 0.87436      | 0.88167      |
| Random Forest  | ohe    | 0.84541      | 0.84997      |
```

	Xgboost		ohe		0.84717		0.85102	
+-----+		+-----+		+-----+		+-----+		+-----+

Observations:

1. One hot encoding features performs better than other encoding technique
2. Linear models (Logistic Regression and SVM) performs better on higher dimension

3 Build Model on frequency encoding feature

3.1 KNN with frequency encoding

```
In [78]: train_df_fc = pd.read_csv('data/train_df_fc.csv')
test_df_fc = pd.read_csv('data/test_df_fc.csv')
```

```
In [79]: train_df_fc.shape, test_df_fc.shape, y_train.shape
```

```
Out[79]: ((32769, 9), (58921, 9), (32769,))
```

```
In [80]: parameters={'n_neighbors':np.arange(1,100, 5)}
clf = RandomizedSearchCV(KNeighborsClassifier(n_jobs=-1),parameters,ran
dom_state=random_state,cv=cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_df_fc,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    5.8s
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    10.4s finished
```


```
In [81]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_n_neighbors')
results
```

Out[81]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	param
0	1.245504	0.191884	0.444588	0.057236	1	{'n_neighbors': 1}
3	0.966955	0.192374	0.425104	0.082150	6	{'n_neighbors': 6}
7	0.915347	0.088786	0.510648	0.062272	16	{'n_neighbors': 16}
5	0.910787	0.067158	0.587300	0.075423	26	{'n_neighbors': 26}
4	0.867675	0.145174	0.647190	0.063128	41	{'n_neighbors': 41}
6	0.877225	0.208047	0.709570	0.113621	56	{'n_neighbors': 56}
2	0.947447	0.128830	0.768725	0.107211	76	{'n_neighbors': 76}
9	0.934347	0.100315	0.597521	0.257969	81	{'n_neighbors': 81}
1	1.011318	0.326486	0.772521	0.073501	86	{'n_neighbors': 86}
8	0.832900	0.196410	0.861454	0.065058	91	{'n_neighbors': 91}

```
In [82]: print_graph(results, 'param_n_neighbors', 'mean_test_score', 'Hyperparameter - No. of neighbors', 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [83]: best_c=best_model.best_params_['n_neighbors']  
best_c
```

Out[83]: 16

```
In [84]: model = KNeighborsClassifier(n_neighbors=best_c,n_jobs=-1)  
model.fit(train_df_fc,y_train)
```

Out[84]: KNeighborsClassifier(n_jobs=-1, n_neighbors=16)

```
In [85]: predictions = model.predict_proba(test_df_fc)[: ,1]  
save_submission(predictions, "knn_fc.csv")
```

[sampleSubmission_knn_fc.csv](#)

12 hours ago by Mayank Gupta

KNN Frequency Encoding

0.79715

0.79125



3.2 SVM with frequency encoding

```
In [86]: C_val = uniform(loc=0, scale=4)
model= LinearSVC(verbose=verbose,random_state=random_state,class_weight
='balanced',max_iter=2000)
parameters={'C':C_val}
clf = RandomizedSearchCV(model,parameters,random_state=random_state,cv=
cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_df_fc,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    17.8s
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    31.9s finished
```

```
[LibLinear]
```

```
In [87]: best_c=best_model.best_params_['C']
best_c
```

```
Out[87]: 3.4647045830997407
```

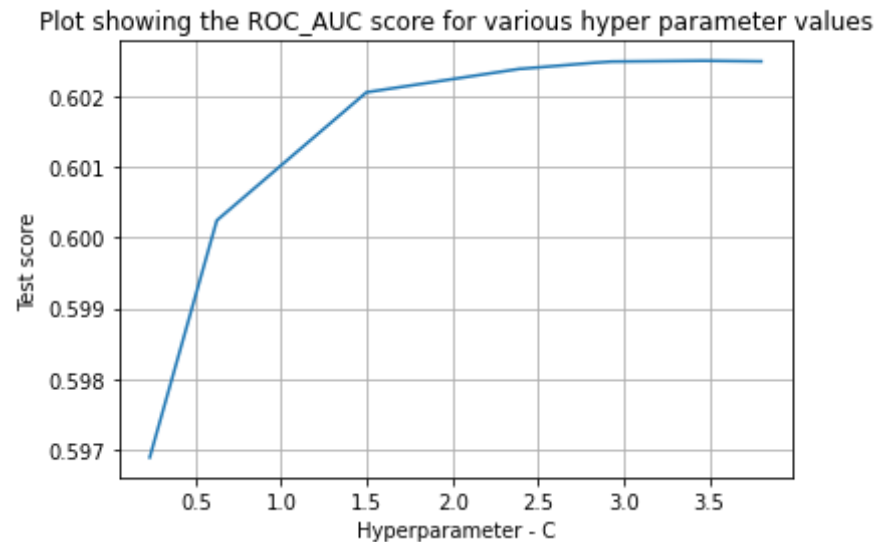
```
In [88]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```

```
Out[88]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
6	0.707697	0.038330	0.009112	0.000330	0.232334	{'C': 0.23233444867279784}
5	1.619894	0.069040	0.009326	0.000816	0.623978	{'C': 0.6239780813448106}
4	1.579539	0.025513	0.008887	0.000250	0.624075	{'C': 0.6240745617697461}
0	3.577643	0.097967	0.009332	0.000580	1.49816	{'C': 1.49816047538945}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
3	5.732613	0.118254	0.009076	0.000258	2.39463	{'C': 2.3946339367881464}
8	5.664490	0.136819	0.009409	0.001280	2.40446	{'C': 2.404460046972835}
9	5.364426	0.630964	0.005462	0.000641	2.83229	{'C': 2.832290311184182}
2	6.995208	0.140802	0.009158	0.000407	2.92798	{'C': 2.9279757672456204}
7	8.412424	0.706659	0.009016	0.000114	3.4647	{'C': 3.4647045830997407}
1	9.626733	0.661829	0.008899	0.000069	3.80286	{'C': 3.8028572256396647}

```
In [89]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C', 'Test score')
```



```
In [90]: #https://stackoverflow.com/questions/26478000/converting-linearsvcs-dec
```

ision-function-to-probabilities-scikit-learn-python

```
model = LinearSVC(C=best_c, verbose=verbose, random_state=random_state, class_weight='balanced', max_iter=2000)
model = CalibratedClassifierCV(model)
model.fit(train_df_fc, y_train)
```

```
[LibLinear][LibLinear][LibLinear][LibLinear][LibLinear]
```

```
Out[90]: CalibratedClassifierCV(base_estimator=LinearSVC(C=3.4647045830997407,
class_weight='balanced',
max_iter=2000, random_state=42,
verbose=2))
```

```
In [91]: predictions = model.predict_proba(test_df_fc)[: ,1]
save_submission(predictions, 'svm_fc.csv')
```

[sampleSubmission_svm_fc.csv](#)

0.60085

0.59550

12 hours ago by [Mayank Gupta](#)

SVM Frequency Encoding

3.3 Logistic Regression with frequency encoding

```
In [92]: C_val = uniform(loc=0, scale=4)
lr= LogisticRegression(verbose=verbose, random_state=random_state, class_weight='balanced', solver='lbfgs', max_iter=500, n_jobs=-1)
parameters={'C':C_val}
clf = RandomizedSearchCV(lr, parameters, random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model = clf.fit(train_df_fc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```



```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    1.7s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed:    8.3s
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed:   20.2s
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   29.2s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed:    0.4s finished
```

```
In [93]: best_c=best_model.best_params_['C']
best_c
```

```
Out[93]: 3.947547746402069
```

```
In [94]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```


```
Out[94]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_
72	0.182546	0.056575	0.010424	0.001924	0.0220885	0.02208846849440959
10	0.230978	0.011034	0.010079	0.002740	0.082338	0.0823379771832097
98	0.246538	0.040240	0.008149	0.000789	0.101677	0.1016765069763807
42	0.264855	0.051649	0.007584	0.001100	0.137554	0.1375540844608739
58	0.294973	0.058159	0.009848	0.001887	0.180909	0.1809091556421522
...
1	0.544588	0.077972	0.010987	0.002171	3.80286	3.802857225639664
34	0.517311	0.047540	0.010529	0.004522	3.86253	3.862528132298237

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_C
50	0.546014	0.078366	0.011045	0.001893	3.87834	3.878338511058234
11	0.487671	0.048548	0.010200	0.002548	3.87964	3.879639408647977
69	0.589011	0.053345	0.010957	0.001444	3.94755	3.94754774640206

100 rows × 14 columns

```
In [95]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C',
, 'Test score')
```



```
In [96]: model = LogisticRegression(C=best_c,verbose=verbose,n_jobs=-1,random_state=random_state,class_weight='balanced',solver='lbfgs')
model.fit(train_df_fc,y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.5s finished
```

```
Out[96]: LogisticRegression(C=3.947547746402069, class_weight='balanced', n_jobs=-1,  
                             random_state=42, verbose=2)
```

```
In [97]: predictions = model.predict_proba(test_df_fc)[: ,1]  
         save_submission(predictions, 'lr_fc.csv')
```

[sampleSubmission_lr_fc.csv](#)

0.59896

0.59778

12 hours ago by [Mayank Gupta](#)

Logistic Regression Frequency Encoding

3.4 Random Forest with frequency encoding

```
In [98]: rfc = RandomForestClassifier(random_state=random_state, class_weight='balanced', n_jobs=-1)  
         clf = RandomizedSearchCV(rfc, get_rf_params(), random_state=random_state,  
                                  cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)  
         best_model = clf.fit(train_df_fc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 18.9s  
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 3.5min  
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed: 8.3min  
/home/auw-mayank/.local/lib/python3.6/site-packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.  
  "timeout or by a memory leak.", UserWarning  
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 12.3min finished
```

```
In [99]: results = pd.DataFrame(best_model.cv_results_)
results.sort_values('mean_test_score',ascending=False,inplace=True)
param_keys=['param_'+str(each) for each in get_rf_params().keys()]
param_keys.append('mean_test_score')
results[param_keys].head(10)
```

Out[99]:

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
78	700	25	2	7	
79	500	25	1	10	
85	1000	20	3	7	
55	200	25	2	5	
27	50	25	2	10	
62	500	25	3	5	
22	200	25	4	10	
6	500	20	2	5	
76	50	25	1	5	
84	1000	25	5	2	

```
In [100]: n_estimators=clf.best_params_['n_estimators']
max_features=clf.best_params_['max_features']
max_depth=clf.best_params_['max_depth']
min_samples_split=clf.best_params_['min_samples_split']
n_estimators,max_features,max_depth,min_samples_split
```


Out[100]: (700, 2, 25, 7)

```
In [101]: model=RandomForestClassifier(n_estimators=n_estimators,max_depth=max_de
pth,max_features=max_features,
min_samples_split=min_samples_split,
random_state=random_state,class_weight='ba
lanced',n_jobs=-1)
```

```
model.fit(train_df_fc,y_train)
```

```
Out[101]: RandomForestClassifier(class_weight='balanced', max_depth=25, max_features=2,  
                                min_samples_split=7, n_estimators=700, n_jobs=-  
                                1,  
                                random_state=42)
```

```
In [103]: features=train_df_fc.columns  
importance=model.feature_importances_  
features=pd.DataFrame({'features':features,'value':importance})  
features=features.sort_values('value',ascending=False)  
sns.barplot('value','features',data=features);  
plt.title('Feature Importance');
```



```
In [106]: predictions = model.predict_proba(test_df_fc)[: ,1]  
save_submission(predictions, 'rf_fc.csv')
```

3.5 Xgboost with frequency encoding

```
In [107]: xgb = XGBClassifier()
clf = RandomizedSearchCV(xgb, get_xgb_params(), random_state=random_state,
cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model=clf.fit(train_df_fc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    8.5s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed:   1.6min
[Parallel(n_jobs=-1)]: Done 349 tasks    | elapsed:   4.4min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   7.3min finished
```

```
In [108]: results = pd.DataFrame(best_model.cv_results_)
results.sort_values('mean_test_score', ascending=False, inplace=True)
param_keys=['param_'+str(each) for each in get_xgb_params().keys()]
param_keys.append('mean_test_score')
results[param_keys].head(10)
```

Out[108]:

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
18	1000	0.048135	0.665922	9	
96	500	0.0979629	0.98664	7	
44	1000	0.060484	0.606429	6	
97	750	0.232385	0.907694	6	
86	1000	0.475848	0.858413	9	

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
53	200	0.540096	0.928319	6	
84	200	0.571989	0.967581	6	
49	500	0.160277	0.393098	8	
62	500	0.0663892	0.328153	9	
14	200	0.374221	0.802197	7	

```
In [109]: colsample_bytree = clf.best_params_['colsample_bytree']
learning_rate=clf.best_params_['learning_rate']
max_depth=clf.best_params_['max_depth']
min_child_weight=clf.best_params_['min_child_weight']
n_estimators=clf.best_params_['n_estimators']
subsample=clf.best_params_['subsample']
colsample_bytree, learning_rate, max_depth, min_child_weight, n_estimators,
subsample
```

```
Out[109]: (0.3308980248526492, 0.04813501017161418, 9, 2, 1000, 0.665922356617496
7)
```

```
In [110]: model = XGBClassifier(colsample_bytree=colsample_bytree, learning_rate=learning_rate,
max_depth=max_depth,
min_child_weight=min_child_weight, n_estimators=n_estimators, subsample=subsample, n_jobs=-1)

model.fit(train_df_fc, y_train)
```

```
Out[110]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=0.3308980248526492,
gamma=0,
gpu_id=-1, importance_type='gain', interaction_constraint
s='',
learning_rate=0.04813501017161418, max_delta_step=0, max_
depth=9,
min_child_weight=2, missing=nan, monotone_constraints
= '()',
```

```


_state=0,
n_estimators=1000, n_jobs=-1, num_parallel_tree=1, random
reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
subsample=0.6659223566174967, tree_method='exact',
validate_parameters=1, verbosity=None)

```

```

In [111]: features=train_df_fc.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features,'value':importance})
features=features.sort_values('value',ascending=False)
sns.barplot('value','features',data=features);
plt.title('Feature Importance');

```



```

In [112]: predictions = model.predict_proba(test_df_fc)[: ,1]
save_submission(predictions, 'xgb_fc.csv')

```

[sampleSubmission_xgb_fc.csv](#)

12 hours ago by [Mayank Gupta](#)

Xgboost Frequency Encoding

0.86987

0.86944

Overview	Data	Notebooks	Discussion	Leaderboard	Rules	Team	My Submissions	Late Submission	
Submission and Description							Private Score	Public Score	Use for Final Score
sampleSubmission_xgb_fc.csv a few seconds ago by Mayank Gupta Xgboost Frequency Encoding							0.86987	0.86944	<input type="checkbox"/>
sampleSubmission_rf_fc.csv a few seconds ago by Mayank Gupta Random Forest Frequency Encoding							0.87299	0.87616	<input type="checkbox"/>
sampleSubmission_lr_fc.csv a minute ago by Mayank Gupta Logistic Regression Frequency Encoding							0.59896	0.59778	<input type="checkbox"/>
sampleSubmission_svm_fc.csv 2 minutes ago by Mayank Gupta SVM Frequency Encoding							0.60085	0.59550	<input type="checkbox"/>
sampleSubmission_knn_fc.csv 2 minutes ago by Mayank Gupta KNN Frequency Encoding							0.79715	0.79125	<input type="checkbox"/>

```
In [113]: from prettytable import PrettyTable

x = PrettyTable(['Model', 'Feature', 'Private Score', 'Public Score'])
x.add_row(['KNN', 'fc', 0.79715, 0.79125])
x.add_row(['SVM', 'fc', 0.60085, 0.59550])
x.add_row(['Logistic Regression', 'fc', 0.59896, 0.59778])
x.add_row(['Random Forest', 'fc', 0.87299, 0.87616])
x.add_row(['Xgboost', 'fc', 0.86987, 0.86944])

print(x)
```

```
+-----+-----+-----+-----+
|      Model      | Feature | Private Score | Public Score |
+-----+-----+-----+-----+
|      KNN        | fc      | 0.79715       | 0.79125      |
|      SVM        | fc      | 0.60085       | 0.5955       |
| Logistic Regression | fc      | 0.59896       | 0.59778      |
```

Random Forest	fc	0.87299	0.87616
Xgboost	fc	0.86987	0.86944

Observations:

1. Tree based models performs better for this feature than linear models
2. KNN is doing good for every feature

4 Build Model using response encoding feature

```
In [114]: train_df_rc = pd.read_csv('data/train_df_rc.csv')
test_df_rc = pd.read_csv('data/test_df_rc.csv')
```

```
In [115]: train_df_rc.shape, test_df_rc.shape, y_train.shape
```

```
Out[115]: ((32769, 9), (58921, 9), (32769,))
```

4.1 KNN with response encoding

```
In [116]: parameters={'n_neighbors':np.arange(1,100, 5)}
clf = RandomizedSearchCV(KNeighborsClassifier(n_jobs=-1),parameters,ran
dom_state=random_state,cv=cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_df_rc,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits


```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    5.2s
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    10.7s finished
```

```
In [117]: results = pd.DataFrame.from_dict(best_model.cv_results_)
          results=results.sort_values('param_n_neighbors')
          results
```

Out[117]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	param
0	0.087703	0.003754	0.270912	0.118324	1	{'n_neighbors': 1}
3	0.161678	0.063581	0.749612	0.191282	6	{'n_neighbors': 6}
7	0.365460	0.149200	1.097595	0.234131	16	{'n_neighbors': 16}
5	0.274667	0.128723	1.106008	0.062541	26	{'n_neighbors': 26}
4	0.196924	0.091312	1.247186	0.043935	41	{'n_neighbors': 41}
6	0.121391	0.037696	1.616742	0.078125	56	{'n_neighbors': 56}
2	0.257782	0.047633	1.718469	0.251959	76	{'n_neighbors': 76}
9	0.259518	0.169725	1.463110	0.571909	81	{'n_neighbors': 81}
1	0.166987	0.110239	2.083340	0.073632	86	{'n_neighbors': 86}
8	0.215055	0.138284	2.578594	0.120702	91	{'n_neighbors': 91}

```
In [118]: print_graph(results, 'param_n_neighbors', 'mean_test_score', 'Hyperparameter - No. of neighbors', 'Test score')
```



```
In [119]: best_c=best_model.best_params_['n_neighbors']  
best_c
```

Out[119]: 91

```
In [120]: model = KNeighborsClassifier(n_neighbors=best_c,n_jobs=-1)  
model.fit(train_df_rc,y_train)
```

Out[120]: KNeighborsClassifier(n_jobs=-1, n_neighbors=91)

```
In [121]: predictions = model.predict_proba(test_df_rc)[: ,1]  
save_submission(predictions, "knn_rc.csv")
```

[sampleSubmission_knn_rc.csv](#)

0.84352

0.85351

12 hours ago by [Mayank Gupta](#)

KNN Response Encoding

4.2 SVM with response encoding

```
In [122]: C_val = uniform(loc=0, scale=4)
model= LinearSVC(verbose=verbose,random_state=random_state,class_weight
='balanced',max_iter=2000)
parameters={'C':C_val}
clf = RandomizedSearchCV(model,parameters,random_state=random_state,cv=
cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_df_rc,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    59.3s
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    1.6min finished
```

```
[LibLinear]
```

```
In [123]: best_c=best_model.best_params_['C']
best_c
```

```
Out[123]: 0.23233444867279784
```

```
In [124]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```


```
Out[124]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
6	3.726326	0.130475	0.009511	0.000833	0.232334	{'C': 0.23233444867279784}
5	9.815611	0.452051	0.009016	0.000450	0.623978	{'C': 0.6239780813448106}
4	9.402443	0.604710	0.009335	0.000597	0.624075	{'C': 0.6240745617697461}
0	20.009992	0.735768	0.009698	0.000943	1.49816	{'C': 1.49816047538945}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
3	19.451480	0.384670	0.013516	0.008621	2.39463	{'C': 2.3946339367881464}
8	17.116975	1.498897	0.008659	0.001653	2.40446	{'C': 2.404460046972835}
9	14.348889	1.305182	0.005841	0.001067	2.83229	{'C': 2.832290311184182}
2	19.717969	0.395682	0.010328	0.002716	2.92798	{'C': 2.9279757672456204}
7	18.970401	0.625358	0.009083	0.000206	3.4647	{'C': 3.4647045830997407}
1	20.206970	0.710743	0.009152	0.000238	3.80286	{'C': 3.8028572256396647}

```
In [125]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C', 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [126]: #https://stackoverflow.com/questions/26478000/converting-linearsvcs-dec
```

ision-function-to-probabilities-scikit-learn-python

```
model = LinearSVC(C=best_c, verbose=verbose, random_state=random_state, class_weight='balanced', max_iter=2000)
model = CalibratedClassifierCV(model)
model.fit(train_df_rc, y_train)
```

```
[LibLinear][LibLinear][LibLinear][LibLinear][LibLinear]
```

```
Out[126]: CalibratedClassifierCV(base_estimator=LinearSVC(C=0.23233444867279784,
class_weight='balanced',
max_iter=2000, random_state=42,
verbose=2))
```

```
In [127]: predictions = model.predict_proba(test_df_rc)[: ,1]
save_submission(predictions, 'svm_rc.csv')
```

[sampleSubmission_svm_rc.csv](#)

0.85160

0.86031

12 hours ago by [Mayank Gupta](#)

SVM Response Encoding

4.3 Logistic Regression with response encoding

```
In [128]: C_val = uniform(loc=0, scale=4)
lr= LogisticRegression(verbose=verbose, random_state=random_state, class_weight='balanced', solver='lbfgs', max_iter=500, n_jobs=-1)
parameters={'C':C_val}
clf = RandomizedSearchCV(lr, parameters, random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model = clf.fit(train_df_rc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    1.8s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed:    9.9s
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed:   24.2s
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   34.9s finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed:    0.5s finished
```

```
In [129]: best_c=best_model.best_params_['C']
best_c
```

```
Out[129]: 3.8783385110582342
```

```
In [130]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```


```
Out[130]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_
72	0.419037	0.018913	0.014351	0.004590	0.0220885	0.02208846849440959
10	0.377327	0.053768	0.012779	0.002567	0.082338	0.0823379771832097
98	0.375209	0.049318	0.010560	0.000751	0.101677	0.1016765069763807
42	0.392769	0.045860	0.014432	0.004220	0.137554	0.1375540844608739
58	0.438569	0.077177	0.013336	0.003642	0.180909	0.1809091556421522
...
1	0.506320	0.115184	0.009774	0.001500	3.80286	3.802857225639664
34	0.435505	0.063518	0.010841	0.001308	3.86253	3.862528132298237

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_C
50	0.560126	0.144761	0.012199	0.003007	3.87834	3.878338511058234
11	0.493065	0.118567	0.013844	0.004191	3.87964	3.879639408647977
69	0.649569	0.209739	0.016016	0.005468	3.94755	3.94754774640206

100 rows × 14 columns

```
In [131]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C',
, 'Test score')
```



```
In [132]: model = LogisticRegression(C=best_c,verbose=verbose,n_jobs=-1,random_state=random_state,class_weight='balanced',solver='lbfgs')
model.fit(train_df_rc,y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 0.4s finished
```

```
Out[132]: LogisticRegression(C=3.8783385110582342, class_weight='balanced', n_jobs=-1,  
                                random_state=42, verbose=2)
```

```
In [133]: predictions = model.predict_proba(test_df_rc)[: ,1]  
          save_submission(predictions, 'lr_rc.csv')
```

[sampleSubmission_lr_rc.csv](#)

0.85322

0.86180

12 hours ago by [Mayank Gupta](#)

LR Response Encoding

4.4 Random Forest with response encoding

```
In [134]: rfc = RandomForestClassifier(random_state=random_state, class_weight='balanced', n_jobs=-1)  
          clf = RandomizedSearchCV(rfc, get_rf_params(), random_state=random_state,  
                                   cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)  
          best_model = clf.fit(train_df_rc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks | elapsed: 19.0s
```

```
[Parallel(n_jobs=-1)]: Done 146 tasks | elapsed: 3.0min
```

```
[Parallel(n_jobs=-1)]: Done 349 tasks | elapsed: 7.0min
```

```
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 10.3min finished
```

```
In [135]: results = pd.DataFrame(best_model.cv_results_)  
          results.sort_values('mean_test_score', ascending=False, inplace=True)  
          param_keys = ['param_' + str(each) for each in get_rf_params().keys()]  
          param_keys.append('mean_test_score')  
          results[param_keys].head(10)
```

Out[135]:

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
68	1000	10	4	20	
26	700	12	4	20	
64	700	10	5	7	
82	700	20	5	20	
41	500	10	3	7	
96	100	10	3	10	
87	200	10	5	2	
11	1000	15	3	7	
85	1000	20	3	7	
25	700	15	4	7	

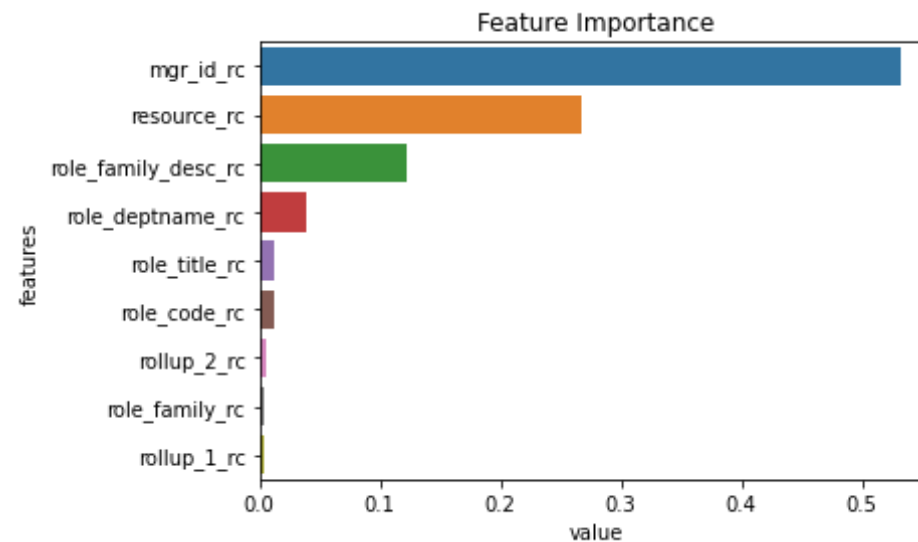
```
In [136]: n_estimators=clf.best_params_['n_estimators']
max_features=clf.best_params_['max_features']
max_depth=clf.best_params_['max_depth']
min_samples_split=clf.best_params_['min_samples_split']
n_estimators,max_features,max_depth,min_samples_split
```

Out[136]: (1000, 4, 10, 20)

```
In [137]: model=RandomForestClassifier(n_estimators=n_estimators,max_depth=max_de
pth,max_features=max_features,
                                     min_samples_split=min_samples_split,
                                     random_state=random_state,class_weight='ba
lanced',n_jobs=-1)
model.fit(train_df_rc,y_train)
```

Out[137]: RandomForestClassifier(class_weight='balanced', max_depth=10, max_features=4,
min_samples_split=20, n_estimators=1000, n_jobs=-1,
random_state=42)

```
In [138]: features=train_df_rc.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features,'value':importance})
features=features.sort_values('value',ascending=False)
sns.barplot('value','features',data=features);
plt.title('Feature Importance');
```



```
In [139]: predictions = model.predict_proba(test_df_rc)[:,:1]
save_submission(predictions, 'rf_rc.csv')
```

[sampleSubmission_rf_rc.csv](#)

0.83136

0.83892

11 hours ago by [Mayank Gupta](#)

RF Response Encoding

4.5 Xgboost with response encoding

```
In [140]: xgb = XGBClassifier()
          clf = RandomizedSearchCV(xgb, get_xgb_params(), random_state=random_state,
          cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
          best_model=clf.fit(train_df_rc, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    7.3s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed:   1.5min
[Parallel(n_jobs=-1)]: Done 349 tasks    | elapsed:   4.0min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   6.7min finished
```

```
In [141]: results = pd.DataFrame(best_model.cv_results_)
          results.sort_values('mean_test_score', ascending=False, inplace=True)
          param_keys=['param_'+str(each) for each in get_xgb_params().keys()]
          param_keys.append('mean_test_score')
          results[param_keys].head(10)
```

Out[141]:

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
1	200	0.0699849	0.601115	5	
28	500	0.0141713	0.222108	5	
7	500	0.017959	0.808397	3	
98	50	0.220131	0.777147	6	
41	20	0.34312	0.996254	6	
58	50	0.454461	0.708911	3	
33	100	0.153737	0.447783	4	
74	20	0.432195	0.763364	6	
94	20	0.591191	0.618218	3	
88	200	0.307937	0.895523	3	

```
In [142]: colsample_bytree = clf.best_params_['colsample_bytree']
learning_rate=clf.best_params_['learning_rate']
max_depth=clf.best_params_['max_depth']
min_child_weight=clf.best_params_['min_child_weight']
n_estimators=clf.best_params_['n_estimators']
subsample=clf.best_params_['subsample']
colsample_bytree, learning_rate, max_depth, min_child_weight, n_estimators,
subsample
```

```
Out[142]: (0.44583275285359114, 0.06998494949080172, 5, 4, 200, 0.601115011743208
8)
```

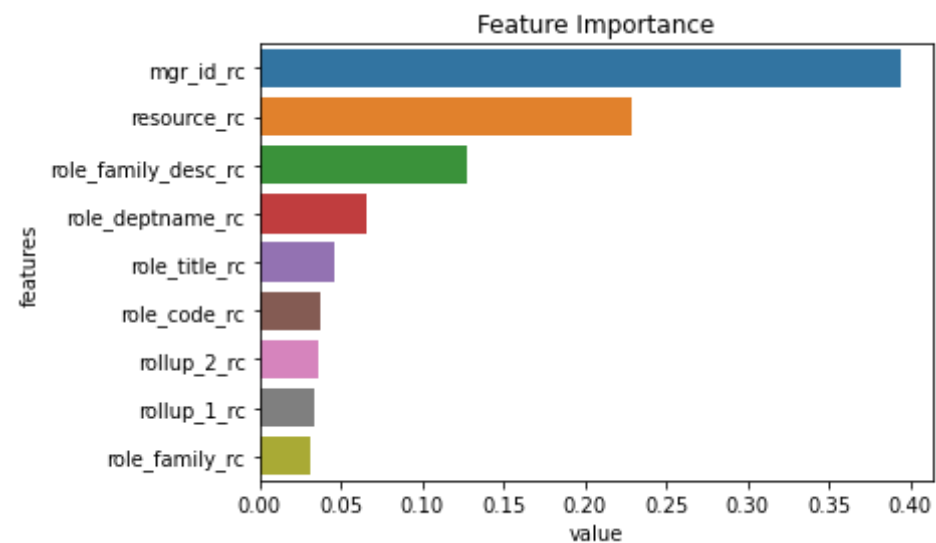
```
In [143]: model = XGBClassifier(colsample_bytree=colsample_bytree, learning_rate=l
earning_rate, max_depth=max_depth,
                                min_child_weight=min_child_weight, n_estimators=n_e
stimators, subsample=subsample, n_jobs=-1)

model.fit(train_df_rc, y_train)
```

```
Out[143]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=0.44583275285359114,
                        gamma=0,
                        gpu_id=-1, importance_type='gain', interaction_constraint
s='',
                        learning_rate=0.06998494949080172, max_delta_step=0, max_
depth=5,
                        min_child_weight=4, missing=nan, monotone_constraints
='()',
                        n_estimators=200, n_jobs=-1, num_parallel_tree=1, random_
state=0,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
                        subsample=0.6011150117432088, tree_method='exact',
                        validate_parameters=1, verbosity=None)
```

```
In [144]: features=train_df_rc.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features, 'value':importance})
features=features.sort_values('value', ascending=False)
```

```
sns.barplot('value', 'features', data=features);  
plt.title('Feature Importance');
```



```
In [145]: predictions = model.predict_proba(test_df_rc)[: ,1]  
save_submission(predictions, 'xgb_rc.csv')
```

[sampleSubmission_xgb_rc.csv](#)

11 hours ago by Mayank Gupta

Xgb Response Encoding

leaderboard

0.84135

0.84190

Search			
Overview	Data	Notebooks	Discussion
Leaderboard	Rules	Team	My Submissions
Late Submission			
Submission and Description	Private Score	Public Score	Use for Final Score
sampleSubmission_xgb_rc.csv just now by Mayank Gupta Xgb Response Encoding	0.84135	0.84190	<input type="checkbox"/>
sampleSubmission_rf_rc.csv a few seconds ago by Mayank Gupta RF Response Encoding	0.83136	0.83892	<input type="checkbox"/>
sampleSubmission_lr_rc.csv 3 minutes ago by Mayank Gupta LR Response Encoding	0.85322	0.86180	<input type="checkbox"/>
sampleSubmission_svm_rc.csv 3 minutes ago by Mayank Gupta SVM Response Encoding	0.85160	0.86031	<input type="checkbox"/>
sampleSubmission_knn_rc.csv 4 minutes ago by Mayank Gupta KNN Response Encoding	0.84352	0.85351	<input type="checkbox"/>

```
In [146]: from prettytable import PrettyTable

x = PrettyTable(['Model', 'Feature', 'Private Score', 'Public Score'])
x.add_row(['KNN', 'rc', 0.84352, 0.85351])
x.add_row(['SVM', 'rc', 0.85160, 0.86031])
x.add_row(['Logistic Regression', 'rc', 0.85322, 0.86180])
x.add_row(['Random Forest', 'rc', 0.83136, 0.83892])
x.add_row(['Xgboost', 'rc', 0.84135, 0.84190])

print(x)
```

Model	Feature	Private Score	Public Score
KNN	rc	0.84352	0.85351
SVM	rc	0.8516	0.86031

Logistic Regression	rc	0.85322	0.8618
Random Forest	rc	0.83136	0.83892
Xgboost	rc	0.84135	0.8419

Observations:

1. Every model performs good for this feature
2. Linear models performs better than Tree based models

5 Build model on SVD feature

```
In [147]: train_svd = pd.read_csv('data/train_svd.csv')
test_svd = pd.read_csv('data/test_svd.csv')
```

```
In [148]: train_svd.shape, test_svd.shape, y_train.shape
```

```
Out[148]: ((32769, 72), (58921, 72), (32769,))
```

5.1 KNN with SVD

```
In [149]: parameters={'n_neighbors':np.arange(1,100, 5)}
clf = RandomizedSearchCV(KNeighborsClassifier(n_jobs=-1),parameters,ran
dom_state=random_state,cv=cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_svd,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed:    19.0s
[Parallel(n_jobs=-1)]: Done 50 out of  50 | elapsed:    38.0s finished
```

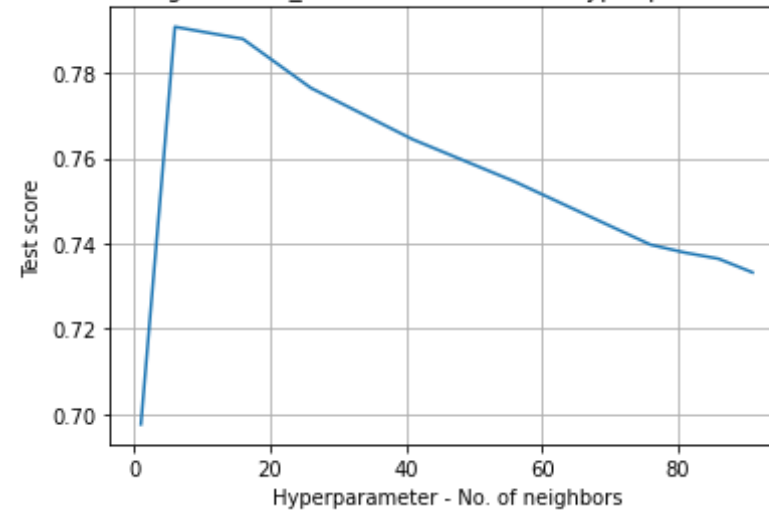
```
In [150]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_n_neighbors')
results
```

Out[150]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	param
0	0.526564	0.011169	1.004034	0.149404	1	{'n_neighbors': 1}
3	1.944336	0.551025	1.655248	0.153292	6	{'n_neighbors': 6}
7	2.016851	0.954735	3.149633	0.351339	16	{'n_neighbors': 16}
5	2.240037	0.818733	3.354217	0.898707	26	{'n_neighbors': 26}
4	1.967969	0.228458	3.715187	1.218701	41	{'n_neighbors': 41}
6	1.281747	0.189357	5.637847	0.125865	56	{'n_neighbors': 56}
2	2.475243	0.578326	5.100802	0.859254	76	{'n_neighbors': 76}
9	2.189955	1.122140	3.499732	1.791585	81	{'n_neighbors': 81}
1	1.204084	0.808660	5.773037	0.319467	86	{'n_neighbors': 86}
8	0.703830	0.153865	7.607220	0.172469	91	{'n_neighbors': 91}

```
In [151]: print_graph(results, 'param_n_neighbors', 'mean_test_score', 'Hyperparameter - No. of neighbors', 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [152]: best_c=best_model.best_params_['n_neighbors']  
best_c
```

```
Out[152]: 6
```

```
In [153]: model = KNeighborsClassifier(n_neighbors=best_c,n_jobs=-1)  
model.fit(train_svd,y_train)
```

```
Out[153]: KNeighborsClassifier(n_jobs=-1, n_neighbors=6)
```

```
In [154]: predictions = model.predict_proba(test_svd)[: ,1]  
save_submission(predictions, "knn_svd.csv")
```

[sampleSubmission_knn_svd.csv](#)

25 minutes ago by [Mayank Gupta](#)

KNN SVD

0.79245

0.78572

5.2 SVM with SVD

```
In [155]: C_val = uniform(loc=0, scale=4)
model= LinearSVC(verbose=verbose,random_state=random_state,class_weight
='balanced',max_iter=2000)
parameters={'C':C_val}
clf = RandomizedSearchCV(model,parameters,random_state=random_state,cv=
cv,verbose=verbose,scoring=scoring,n_jobs=-1)
best_model = clf.fit(train_svd,y_train)
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 1.5min
[Parallel(n_jobs=-1)]: Done 50 out of 50 | elapsed: 2.7min finished
```

```
[LibLinear]
```

```
In [156]: best_c=best_model.best_params_['C']
best_c
```

```
Out[156]: 3.8028572256396647
```


```
In [157]: results = pd.DataFrame.from_dict(best_model.cv_results_)
results=results.sort_values('param_C')
results
```

```
Out[157]:
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
6	2.968880	0.150177	0.016558	0.002776	0.232334	{'C': 0.232334444867279784}
5	6.587434	0.301619	0.012963	0.000876	0.623978	{'C': 0.6239780813448106}
4	6.343227	0.065596	0.012513	0.000442	0.624075	{'C': 0.6240745617697461}
0	17.291072	1.080804	0.016185	0.007285	1.49816	{'C': 1.49816047538945}

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params
3	29.455153	0.767611	0.015290	0.003827	2.39463	{'C': 2.3946339367881464}
8	29.985117	1.952692	0.012303	0.002009	2.40446	{'C': 2.404460046972835}
9	27.816009	2.354661	0.007257	0.000748	2.83229	{'C': 2.832290311184182}
2	36.411890	2.210566	0.012634	0.000489	2.92798	{'C': 2.9279757672456204}
7	44.704208	1.334583	0.012921	0.000702	3.4647	{'C': 3.4647045830997407}
1	48.909459	2.384907	0.016626	0.005413	3.80286	{'C': 3.8028572256396647}

```
In [158]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C', 'Test score')
```



```
In [159]: #https://stackoverflow.com/questions/26478000/converting-linearsvcs-dec
```

ision-function-to-probabilities-sci-kit-learn-python

```
model = LinearSVC(C=best_c, verbose=verbose, random_state=random_state, class_weight='balanced', max_iter=2000)
model = CalibratedClassifierCV(model)
model.fit(train_svd, y_train)
```

```
[LibLinear][LibLinear][LibLinear][LibLinear][LibLinear]
```

```
Out[159]: CalibratedClassifierCV(base_estimator=LinearSVC(C=3.8028572256396647,
class_weight='balanced',
max_iter=2000, random_state=42,
verbose=2))
```

```
In [160]: predictions = model.predict_proba(test_svd)[: ,1]
save_submission(predictions, 'svm_svd.csv')
```

[sampleSubmission_svm_svd.csv](#)

0.63648

0.63806

24 minutes ago by [Mayank Gupta](#)

SVM SVD

5.3 Logistic Regression with SVD

```
In [161]: C_val = uniform(loc=0, scale=4)
lr= LogisticRegression(verbose=verbose, random_state=random_state, class_weight='balanced', solver='lbfgs', max_iter=500, n_jobs=-1)
parameters={'C':C_val}
clf = RandomizedSearchCV(lr, parameters, random_state=random_state, cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
best_model = clf.fit(train_svd, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

```
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 18.9s
[Parallel(n_jobs=-1)]: Done 146 tasks     | elapsed: 1.4min
[Parallel(n_jobs=-1)]: Done 349 tasks     | elapsed: 3.4min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 5.1min finished
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 14.7s finished
```

```
In [162]: best_c=best_model.best_params_['C']
          best_c
```

```
Out[162]: 3.947547746402069
```

```
In [163]: results = pd.DataFrame.from_dict(best_model.cv_results_)
          results=results.sort_values('param_C')
          results
```

```
Out[163]:
```


	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_
72	1.042337	0.102578	0.035350	0.017093	0.0220885	0.02208846849440959
10	1.598916	0.145812	0.029739	0.011829	0.082338	0.0823379771832097
98	1.946853	0.134926	0.031715	0.005840	0.101677	0.1016765069763807
42	1.629183	0.083901	0.022801	0.003346	0.137554	0.1375540844608739
58	1.859177	0.136240	0.025097	0.007581	0.180909	0.1809091556421522
...
1	7.015733	0.438911	0.026228	0.005973	3.80286	3.802857225639664
34	6.146194	0.345854	0.028273	0.008512	3.86253	3.862528132298237

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	param_C
50	6.477816	0.289469	0.026133	0.013208	3.87834	3.878338511058234
11	6.560687	0.689193	0.033242	0.007402	3.87964	3.879639408647977
69	7.289347	0.445322	0.030326	0.003518	3.94755	3.94754774640206

100 rows × 14 columns

```
In [164]: print_graph(results, 'param_C', 'mean_test_score', 'Hyperparameter - C',
, 'Test score')
```

Plot showing the ROC_AUC score for various hyper parameter values



```
In [165]: model = LogisticRegression(C=best_c,verbose=verbose,n_jobs=-1,random_state=random_state,class_weight='balanced',solver='lbfgs')
model.fit(train_svd,y_train)
```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.


```
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 3.5s finished
```

```
Out[165]: LogisticRegression(C=3.947547746402069, class_weight='balanced', n_jobs=-1,  
                             random_state=42, verbose=2)
```

```
In [166]: predictions = model.predict_proba(test_svd)[: ,1]  
          save_submission(predictions, 'lr_svd.csv')
```

[sampleSubmission_lr_svd.csv](#)

0.63255

0.63314

24 minutes ago by [Mayank Gupta](#)

Logistic Regression SVD

5.4 Random Forest with SVD

```
In [167]: rfc = RandomForestClassifier(random_state=random_state, class_weight='balanced', n_jobs=-1)  
          clf = RandomizedSearchCV(rfc, get_rf_params(), random_state=random_state,  
                                   cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)  
          best_model = clf.fit(train_svd, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=-1)]: Done 25 tasks | elapsed: 36.2s  
/home/auw-mayank/.local/lib/python3.6/site-packages/joblib/externals/loky/process_executor.py:691: UserWarning: A worker stopped while some jobs were given to the executor. This can be caused by a too short worker timeout or by a memory leak.
```

"timeout or by a memory leak.", UserWarning

```
[Parallel(n_jobs=-1)]: Done 146 tasks | elapsed: 6.8min
```

```
[Parallel(n_jobs=-1)]: Done 349 tasks | elapsed: 16.5min
```

```
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 24.5min finished
```

```
In [168]: results = pd.DataFrame(best_model.cv_results_)
results.sort_values('mean_test_score',ascending=False,inplace=True)
param_keys=['param_'+str(each) for each in get_rf_params().keys()]
param_keys.append('mean_test_score')
results[param_keys].head(10)
```

Out[168]:

	param_n_estimators	param_max_depth	param_max_features	param_min_samples_split	mean
84	1000	25	5	2	
20	1000	25	3	2	
33	700	25	4	2	
22	200	25	4	10	
78	700	25	2	7	
85	1000	20	3	7	
62	500	25	3	5	
82	700	20	5	20	
79	500	25	1	10	
92	500	20	3	2	

```
In [169]: n_estimators=clf.best_params_['n_estimators']
max_features=clf.best_params_['max_features']
max_depth=clf.best_params_['max_depth']
min_samples_split=clf.best_params_['min_samples_split']
n_estimators,max_features,max_depth,min_samples_split
```


Out[169]: (1000, 5, 25, 2)

```
In [170]: model=RandomForestClassifier(n_estimators=n_estimators,max_depth=max_de
pth,max_features=max_features,
min_samples_split=min_samples_split,
random_state=random_state,class_weight='ba
lanced',n_jobs=-1)
```

```
model.fit(train_svd,y_train)
```

```
Out[170]: RandomForestClassifier(class_weight='balanced', max_depth=25, max_features=5,
                                n_estimators=1000, n_jobs=-1, random_state=42)
```

```
In [171]: features=train_svd.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features,'value':importance})
features=features.sort_values('value',ascending=False)
sns.barplot('value','features',data=features);
plt.title('Feature Importance');
```



```
In [172]: predictions = model.predict_proba(test_svd)[: ,1]
          save_submission(predictions, 'rf_svd.csv')
```

sampleSubmission_rf_svd.csv

18 minutes ago by [Mayank Gupta](#)

Random Forest SVD

0.87119

0.86924

5.5 Xgboost with SVD

```
In [173]: xgb = XGBClassifier()
          clf = RandomizedSearchCV(xgb, get_xgb_params(), random_state=random_state,
          cv=cv, verbose=verbose, n_iter=100, scoring=scoring, n_jobs=-1)
          best_model=clf.fit(train_svd, y_train)
```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
/home/auw-mayank/.local/lib/python3.6/site-packages/joblib/externals/lo
ky/process_executor.py:691: UserWarning: A worker stopped while some jo
bs were given to the executor. This can be caused by a too short worker
timeout or by a memory leak.
  "timeout or by a memory leak.", UserWarning
[Parallel(n_jobs=-1)]: Done 25 tasks      | elapsed: 54.6s
[Parallel(n_jobs=-1)]: Done 146 tasks    | elapsed: 9.9min
[Parallel(n_jobs=-1)]: Done 349 tasks    | elapsed: 28.0min
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 50.8min finished
```

```
In [174]: results = pd.DataFrame(best_model.cv_results_)
          results.sort_values('mean_test_score', ascending=False, inplace=True)
          param_keys=['param_'+str(each) for each in get_xgb_params().keys()]
          param_keys.append('mean_test_score')
          results[param_keys].head(10)
```

Out[174]:

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
62	500	0.0663892	0.328153	9	
18	1000	0.048135	0.665922	9	
96	500	0.0979629	0.98664	7	
44	1000	0.060484	0.606429	6	
8	750	0.0686033	0.683264	6	

	param_n_estimators	param_learning_rate	param_subsample	param_max_depth	param_colsan
97	750	0.232385	0.907694	6	
49	500	0.160277	0.393098	8	
80	1000	0.385564	0.905351	3	
53	200	0.540096	0.928319	6	
78	1000	0.576551	0.94023	6	

```
In [175]: colsample_bytree = clf.best_params_['colsample_bytree']
learning_rate=clf.best_params_['learning_rate']
max_depth=clf.best_params_['max_depth']
min_child_weight=clf.best_params_['min_child_weight']
n_estimators=clf.best_params_['n_estimators']
subsample=clf.best_params_['subsample']
colsample_bytree, learning_rate, max_depth, min_child_weight, n_estimators,
subsample
```

```
Out[175]: (0.375582952639944, 0.06638916390452139, 9, 3, 500, 0.3281526674747319
3)
```

```
In [176]: model = XGBClassifier(colsample_bytree=colsample_bytree, learning_rate=learning_rate,
max_depth=max_depth,
min_child_weight=min_child_weight, n_estimators=n_estimators, subsample=subsample, n_jobs=-1)

model.fit(train_svd, y_train)
```

```
Out[176]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=0.375582952639944, gamma=0,
gpu_id=-1, importance_type='gain', interaction_constraint_s='',
learning_rate=0.06638916390452139, max_delta_step=0, max_depth=9,
min_child_weight=3, missing=nan, monotone_constraints='()',
```

```


state=0,
n_estimators=500, n_jobs=-1, num_parallel_tree=1, random_
reg_alpha=0, reg_lambda=1, scale_pos_weight=1,
subsample=0.32815266747473193, tree_method='exact',
validate_parameters=1, verbosity=None)

```

```

In [177]: features=train_svd.columns
importance=model.feature_importances_
features=pd.DataFrame({'features':features,'value':importance})
features=features.sort_values('value',ascending=False)
sns.barplot('value','features',data=features);
plt.title('Feature Importance');

```



```

In [178]: predictions = model.predict_proba(test_svd)[:,:1]
save_submission(predictions, 'xgb_svd.csv')

```

[sampleSubmission_xgb_svd.csv](#)

18 minutes ago by [Mayank Gupta](#)

Xgboost SVD

0.86909

0.86664

Overview	Data	Notebooks	Discussion	Leaderboard	Rules	Team	My Submissions	Late Submission
Submission and Description					Private Score	Public Score	Use for Final Score	
sampleSubmission_xgb_svd.csv just now by Mayank Gupta Xgboost SVD					0.86909	0.86664	<input type="checkbox"/>	
sampleSubmission_rf_svd.csv a few seconds ago by Mayank Gupta Random Forest SVD					0.87119	0.86924	<input type="checkbox"/>	
sampleSubmission_lr_svd.csv 6 minutes ago by Mayank Gupta Logistic Regression SVD					0.63255	0.63314	<input type="checkbox"/>	
sampleSubmission_svm_svd.csv 7 minutes ago by Mayank Gupta SVM SVD					0.63648	0.63806	<input type="checkbox"/>	
sampleSubmission_knn_svd.csv 7 minutes ago by Mayank Gupta KNN SVD					0.79245	0.78572	<input type="checkbox"/>	

```
In [179]: from prettytable import PrettyTable

x = PrettyTable(['Model', 'Feature', 'Private Score', 'Public Score'])
x.add_row(['KNN', 'svd', 0.79245, 0.78572])
x.add_row(['SVM', 'svd', 0.63648, 0.63806])
x.add_row(['Logistic Regression', 'svd', 0.63255, 0.63314])
x.add_row(['Random Forest', 'svd', 0.87119, 0.86924])
x.add_row(['Xgboost', 'svd', 0.86909, 0.86664])

print(x)
```

Model	Feature	Private Score	Public Score
KNN	svd	0.79245	0.78572
SVM	svd	0.63648	0.63806
Logistic Regression	svd	0.63255	0.63314

	Random Forest		svd		0.87119		0.86924	
	Xgboost		svd		0.86909		0.86664	
+-----+		+-----+		+-----+		+-----+		+-----+

Observations:

1. Tree based models works better than linear model
2. KNN is performing overall good

We have to improve our model to reach into 5-10% on kaggle

In [180]: `# https://www.kaggle.com/mitribunskiy/tutorial-catboost-overview`

In [181]: `# https://www.kaggle.com/prashant111/catboost-classifier-tutorial`

<https://catboost.ai/>

CatBoost is a high-performance open source library for gradient boosting on decision trees

About

CatBoost is an algorithm for gradient boosting on decision trees. It is developed by Yandex researchers and engineers, and is used for search, recommendation systems, personal assistant, self-driving cars, weather prediction and many other tasks at Yandex and in other companies, including CERN, Cloudflare, Careem taxi. It is in open-source and can be used by anyone.

Features

1. Reduce time spent on parameter tuning, because CatBoost provides great results with default parameters
2. Improve your training results with CatBoost that allows you to use non-numeric factors, instead of having to pre-process your data or spend time and effort turning it to numbers.
3. Reduce overfitting when constructing your models with a novel gradient-boosting scheme.
4. Apply your trained model quickly and efficiently even to latency-critical tasks using CatBoost's model applier

```
In [182]: params = {  
            'loss_function': 'Logloss',  
            'eval_metric': 'AUC',  
            'cat_features': list(range(train_data.shape[1])),  
            'verbose': 100,  
            'random_seed': random_state  
        }
```

```
In [183]: clf= CatBoostClassifier(**params)  
          clf.fit(train_data,y_train)  
  
Learning rate set to 0.045713  
0:      total: 99.2ms    remaining: 1m 39s  
100:    total: 2.33s    remaining: 20.7s  
200:    total: 5.6s     remaining: 22.3s  
300:    total: 8.79s    remaining: 20.4s  
400:    total: 11.9s    remaining: 17.8s  
500:    total: 15.2s    remaining: 15.2s  
600:    total: 18.4s    remaining: 12.2s  
700:    total: 21.7s    remaining: 9.27s  
800:    total: 24.8s    remaining: 6.16s  
900:    total: 28s      remaining: 3.08s  
999:    total: 31.1s    remaining: 0us
```

```
Out[183]: <catboost.core.CatBoostClassifier at 0x7f361d4d8780>
```

```
In [184]: predictions = clf.predict_proba(test_data)[: ,1]
```

```
In [185]: save_submission(predictions, 'catboost.csv')
```

Submission and Description	Private Score	Public Score	Use for Final Score
sampleSubmission_catboost.csv a minute ago by Mayank Gupta Finally I am using catboost model for Amazon Employee Access.	0.90889	0.91483	<input type="checkbox"/>

**Catboost perform better than all our previous models
and it's AUC score is much better than previous models
so I am selecting this for predicting future data**

```
In [186]: # Save model on disk  
pickle.dump(clf, open('models/catboost_model.pkl', 'wb'))
```

```
In [ ]:
```